

Essential R Skills

UMN-Morris Statistics Discipline

2023-06-21

Contents

1	Motivation	5
2	Getting Started	7
2.1	Packages	7
2.2	The tidyverse Package	7
2.3	Gapminder Data:	7
2.4	Set Working Directory	8
2.5	Reading Data From a CSV File	8
3	Overview of a Dataframe	11
3.1	glimpse	11
3.2	dim	11
3.3	head	12
3.4	tail	12
3.5	summary	13
3.6	Dataframe Details: funModeling package	13
3.7	Dataframe Details: skimr package	14
3.8	describe: Hmisc package	15
4	Introduction to Data Wrangling	17
4.1	Tidy Data	17
4.2	Subset Observations using filter	17
4.3	Subset using multiple conditions	18
4.4	Saving as a new dataframe	19
4.5	Subset using top_n	20
4.6	Subset using select	20
4.7	Order using arrange	22
4.8	Grouped Filter	23
4.9	New Variables Using mutate	24
4.10	Simple Counting Using tally() and count()	25
4.11	Missing Values	27
5	Univariate Graphical Displays	29
5.1	Overview of ggplot	29
5.2	A Quantitative Variable	29
5.3	Displays of a Categorical Variable	43
6	Summary Statistics For One Variable	51
6.1	One Quantitative Variable	51
6.2	One Categorical Variable	55
7	Exploratory Data Analysis For One Quantitative Variable: by Groups	59
7.1	Summary Statistics: dplyr	59
7.2	Summary Statistics: gtsummary	61

7.3	Summary Statistics: skimr	61
7.4	Graphical Displays of a quantitative variable, separated by groups	61
8	Analysis of One Categorical Variable by another categorical variable	71
8.1	Tables	71
8.2	Graphical Displays	73

Chapter 1

Motivation

We have found that students enter our courses with wide variation in experience and comfort using statistical software for computation and making graphical displays. This document represents our expectations for the basic R skills that students should know upon completing an introductory course in statistics. Analysis methods may appear at times in this document, but the emphasis here is upon basic R usage for data wrangling, and exploratory data analysis using numeric and graphical methods.

Chapter 2

Getting Started

2.1 Packages

When you start R studio, basic functionality is initially available. However, in most projects we will want to use some special code and functions contained in packages that are not initially available when R starts. Before packages can be used in our analyses, they must be installed in our R workspace. We presume that the R Studio development environment is being used by our students. Any package can be installed by clicking the “Packages” tab in the lower right panel of the R Studio workspace. Then click “Install” to produce an entry bar where you type the name of the desired package. You can also instruct R Studio to install a package directly by typing “install.packages(“tidyverse”)” in the “Console” window in the lower left area of R Studio. This command will execute and install the *tidyverse* package.

Once a package is installed into your R Studio environment, you make it available by loading with the `library()` command. For this document, some additional packages are needed, and are loaded in the next code block. The *knitr* and *tidyverse* packages have been previously installed. If you attempt to load a package (in this case *zelig*) that has not been installed, you will get an error message similar to this:

```
Error in library(zelig) : there is no package called ‘zelig’
```

```
# this block loads R packages that may be needed for the analysis.  
library(knitr)  
library(tidyverse)
```

2.2 The tidyverse Package

The tidyverse package is very special - it is a package of other packages. The tidyverse website tidyverse describes the tidyverse as: *The tidyverse is an opinionated collection of R packages designed for data science. All packages share an underlying design philosophy, grammar, and data structures.*

The most important packages inside the tidyverse package for this document are: *dplyr*, *magrittr*, and *ggplot2*.

2.3 Gapminder Data:

This dataset (named *gapminder*) is contained in an R package called *gapminder*, and needs to be loaded before the dataset can be used.

```
library(gapminder)
```

2.4 Set Working Directory

In the “Files” tab in the lower right portion of the R Studio work area, you can choose where you want to store files and conduct your work by navigating to a suitable folder by making folders and sub-folders and then clicking to navigate to a suitable work area.

You should notify R Studio and the R software to this location called the working directory. Once you have navigated to where you want files, data, and results to reside, you notify R Studio by clicking the blue “More” gear and choose “set working directory.” This will help R understand where to expect files and data to be located.

One of the most common problems students experience is that they work on files in a location not specified as the “working directory.”

2.5 Reading Data From a CSV File

The most common way to read data into R is from an excel spreadsheet that has been saved into a comma-separated-values (csv) file. This means that data elements are separated from each other by commas “,”.

2.5.1 CSV File Located in Working Directory

We consider a data file named (file.csv) that contains variable names in the first row of the file. Place this file in your working directory and read,

```
dataframe <- read.csv("file.csv",header=TRUE)
```

A frequent issue with `read.csv` is that character variables are automatically converted to factor/categorical variables. This may not be a good choice in many instances. To gain full control of how this is handled, you can prevent this kind of auto-conversion by using the `stringsAsFactors` option.

```
cardata <- read.csv(file = 'carspeeds.csv', stringsAsFactors = FALSE)
```

The `readr` package inside the *tidyverse* family of packages has a slightly nicer read function you should know about. We use the `readr::` prefix to inform readers that the `read_csv` function resides in the `readr` package. This read function will not auto-convert character variables to category/factor variables.

```
dataframe <- readr::read_csv("file.csv",col_names = TRUE)
```

2.5.2 Read Data From a URL Web Location

If data is located in an accessible location on the internet, you can simply add the location pathway or URL address to the `read.csv` command to access the data.

```
#import data from URL
data <- read.csv('https://raw.githubusercontent.com/Statology/Miscellaneous/main/basketball_data.csv')

#view first five rows
head(data)
```

```
##   player assists points
## 1      A        6     12
## 2      B        7     19
## 3      C       14      7
## 4      D        4      6
## 5      E        5     10
```


Here is a similar command using the *readr* package.

```
#import data from URL
data3 <- data.table::fread('https://raw.githubusercontent.com/Statology/Miscellaneous/main/basketball_data')

#view first five rows
head(data3)
```

```
##      player assists points
## 1:      A         6      12
## 2:      B         7      19
## 3:      C        14         7
## 4:      D         4         6
## 5:      E         5        10
```

2.5.3 Reading Data From Excel Files

Reading data directly from excel spreadsheets is more complex and you should read documentation for the *readxl* package. Here we illustrate the basic idea.

```
library(readxl)
# Read sheet names 'sheet_name' from excel file
df = read_excel("/Users/new_file.xlsx", sheet='sheet_name')
#or
# Read sheet 3 from excel file
df = read_excel("/Users/new_file.xlsx", sheet=3)
```


Chapter 3

Overview of a Dataframe

Datasets in R are usually called dataframes or tibbles. The distinction between these names is not important for our purposes - we will usually refer to a dataset as a dataframe.

3.1 glimpse

Let's look at what is inside the `gapminder` dataset using the `glimpse` command from the `dplyr` package. The `dplyr` package is contained in the package “tidyverse” that was loaded previously. The `glimpse(gapminder)` command would have executed without any errors. We use the `dplyr::` prefix to inform readers that the `glimpse` function resides in the `dplyr` package.

```
# the next command would also execute if  
# dplyr or tidyverse was previously loaded with library(dplyr)  
#glimpse(gapminder)  
# instead, we note the package (dplyr) where command (glimpse) is found:  
dplyr::glimpse(gapminder)
```

```
## Rows: 1,704  
## Columns: 6  
## $ country   <fct> "Afghanistan", "Afghanistan", "Afghanistan", "Afghanistan", "Afghanistan", "Afghanist  
## $ continent <fct> Asia, Asia, Asia, Asia, Asia, Asia, Asia, Asia, Asia, Asia, Asia, Asia, Asia, Europe, Europ  
## $ year      <int> 1952, 1957, 1962, 1967, 1972, 1977, 1982, 1987, 1992, 1997, 2002, 2007, 1952, 1957, 1  
## $ lifeExp   <dbl> 28.801, 30.332, 31.997, 34.020, 36.088, 38.438, 39.854, 40.822, 41.674, 41.763, 42.12  
## $ pop       <int> 8425333, 9240934, 10267083, 11537966, 13079460, 14880372, 12881816, 13867957, 1631792  
## $ gdpPercap <dbl> 779.4453, 820.8530, 853.1007, 836.1971, 739.9811, 786.1134, 978.0114, 852.3959, 649.3
```

This shows it contains economic and demographic information about different countries across years. There are 1704 rows (observations) and 6 columns (variables).

Each variable name is listed along with a variable type designation.

- `fct`: means a factor variable, also known as a categorical variable.
- `int`: means a quantitative variable that takes only integer or whole number values.
- `dbl`: means double precision, a quantitative variable that is essentially continuous - taking decimal values.

3.2 dim

The `dim` command quickly tells us the number of rows (observations) and columns (variables) in a data frame.

```
dim(gapminder)
```

```
## [1] 1704    6
```

3.3 head

By default, the `head` command will show the first 6 rows of the dataset `gapminder`. Datasets in R are called “dataframes.” The `gapminder` dataframe is denoted as a “tibble” which is a type of dataframe.

Options to the `head` command can change the rows displayed.

```
# default is to show 6 rows
```

```
head(gapminder)
```

```
## # A tibble: 6 x 6
##   country    continent  year lifeExp      pop gdpPercap
##   <fct>      <fct>    <int>  <dbl>    <int>    <dbl>
## 1 Afghanistan Asia      1952   28.8  8425333    779.
## 2 Afghanistan Asia      1957   30.3  9240934    821.
## 3 Afghanistan Asia      1962   32.0 10267083    853.
## 4 Afghanistan Asia      1967   34.0 11537966    836.
## 5 Afghanistan Asia      1972   36.1 13079460    740.
## 6 Afghanistan Asia      1977   38.4 14880372    786.
```

```
# show only 4 rows...
```

```
head(gapminder,n=4)
```

```
## # A tibble: 4 x 6
##   country    continent  year lifeExp      pop gdpPercap
##   <fct>      <fct>    <int>  <dbl>    <int>    <dbl>
## 1 Afghanistan Asia      1952   28.8  8425333    779.
## 2 Afghanistan Asia      1957   30.3  9240934    821.
## 3 Afghanistan Asia      1962   32.0 10267083    853.
## 4 Afghanistan Asia      1967   34.0 11537966    836.
```

3.4 tail

By default, the `tail` command will show the last 6 rows of the dataset `gapminder`.

Options to the `tail` command can change the rows displayed.

```
# default is to show 6 rows
```

```
tail(gapminder)
```

```
## # A tibble: 6 x 6
##   country    continent  year lifeExp      pop gdpPercap
##   <fct>      <fct>    <int>  <dbl>    <int>    <dbl>
## 1 Zimbabwe Africa      1982   60.4  7636524    789.
## 2 Zimbabwe Africa      1987   62.4  9216418    706.
## 3 Zimbabwe Africa      1992   60.4 10704340    693.
## 4 Zimbabwe Africa      1997   46.8 11404948    792.
## 5 Zimbabwe Africa      2002   40.0 11926563    672.
## 6 Zimbabwe Africa      2007   43.5 12311143    470.
```

```
# show only 4 rows...
```

```
tail(gapminder,n=4)
```

```
## # A tibble: 4 x 6
##   country continent year lifeExp      pop gdpPercap
##   <fct>    <fct>    <int>   <dbl>   <int>   <dbl>
## 1 Zimbabwe Africa     1992    60.4 10704340    693.
## 2 Zimbabwe Africa     1997    46.8 11404948    792.
## 3 Zimbabwe Africa     2002    40.0 11926563    672.
## 4 Zimbabwe Africa     2007    43.5 12311143    470.
```

3.5 summary

This command shows a basic summary of the values in each variable.

```
# A basic, base R command
summary(gapminder)
```

```
##           country      continent      year      lifeExp      pop      gdpPercap
## Afghanistan: 12 Africa :624 Min. :1952 Min. :23.60 Min. :6.001e+04 Min. : 241.2
## Albania : 12 Americas:300 1st Qu.:1966 1st Qu.:48.20 1st Qu.:2.794e+06 1st Qu.: 1202.1
## Algeria : 12 Asia :396 Median :1980 Median :60.71 Median :7.024e+06 Median : 3531.8
## Angola : 12 Europe :360 Mean :1980 Mean :59.47 Mean :2.960e+07 Mean : 7215.3
## Argentina : 12 Oceania : 24 3rd Qu.:1993 3rd Qu.:70.85 3rd Qu.:1.959e+07 3rd Qu.: 9325.5
## Australia : 12 Max. :2007 Max. :82.60 Max. :1.319e+09 Max. :113523.1
## (Other) :1632
```

The next command illustrates a “pipe” - here the dataframe `gapminder` is “piped” into the `summary` function to be processed. Note the same output is produced as using `summary(gapminder)`. Note, the pipe operation `%>%` is contained in tidyverse package: *magrittr* which is loaded when *tidyverse* is loaded.

```
# Same idea, as summary(gapminder) but using tidyverse pipe
gapminder %>% summary()
```

```
##           country      continent      year      lifeExp      pop      gdpPercap
## Afghanistan: 12 Africa :624 Min. :1952 Min. :23.60 Min. :6.001e+04 Min. : 241.2
## Albania : 12 Americas:300 1st Qu.:1966 1st Qu.:48.20 1st Qu.:2.794e+06 1st Qu.: 1202.1
## Algeria : 12 Asia :396 Median :1980 Median :60.71 Median :7.024e+06 Median : 3531.8
## Angola : 12 Europe :360 Mean :1980 Mean :59.47 Mean :2.960e+07 Mean : 7215.3
## Argentina : 12 Oceania : 24 3rd Qu.:1993 3rd Qu.:70.85 3rd Qu.:1.959e+07 3rd Qu.: 9325.5
## Australia : 12 Max. :2007 Max. :82.60 Max. :1.319e+09 Max. :113523.1
## (Other) :1632
```

3.6 Dataframe Details: funModeling package

The *funModeling* package contains the `df_status` command which also summarizes a dataframe - showing different aspects like missing values, percentage of zero values, and also the number of unique values.

```
funModeling::df_status(gapminder)
```

```
##   variable q_zeros p_zeros q_na p_na q_inf p_inf      type unique
## 1 country      0      0    0    0    0    0 factor     142
## 2 continent    0      0    0    0    0    0 factor       5
## 3 year         0      0    0    0    0    0 integer     12
## 4 lifeExp      0      0    0    0    0    0 numeric    1626
## 5 pop          0      0    0    0    0    0 integer    1704
## 6 gdpPercap    0      0    0    0    0    0 numeric    1704
```

```

di=funModeling::data_integrity(gapminder)
# returns a detailed summary of all variables
print(di)

## $vars_num_with_NA
## [1] variable q_na      p_na
## <0 rows> (or 0-length row.names)
##
## $vars_cat_with_NA
## [1] variable q_na      p_na
## <0 rows> (or 0-length row.names)
##
## $vars_cat_high_card
##           variable unique
## country  country      142
##
## $MAX_UNIQUE
## [1] 35
##
## $vars_one_value
## character(0)
##
## $vars_cat
## [1] "country"      "continent"
##
## $vars_num
## [1] "year"         "lifeExp"      "pop"          "gdpPercap"
##
## $vars_char
## character(0)
##
## $vars_factor
## [1] "country"      "continent"
##
## $vars_other
## character(0)

```

```
## -----
## Group variables          None
##
## -- Variable type: factor -----
##   skim_variable n_missing complete_rate ordered n_unique top_counts
## 1 country        0              1 FALSE      142 Afg: 12, Alb: 12, Alg: 12, Ang: 12
## 2 continent      0              1 FALSE        5 Afr: 624, Asi: 396, Eur: 360, Ame: 300
##
## -- Variable type: numeric -----
##   skim_variable n_missing complete_rate      mean      sd      p0      p25      p50      p75
## 1 year          0              1    1980.    17.3  1952    1966.    1980.    1993.
## 2 lifeExp       0              1     59.5    12.9  23.6    48.2    60.7    70.8
## 3 pop           0              1 29601212. 106157897. 60011 2793664 7023596. 19585222. 1
## 4 gdpPercap     0              1    7215.    9857.   241.   1202.   3532.   9325.
```

3.8 describe: Hmisc package

The *Hmisc* package contains the `describe` function that gives a helpful overview of numeric and categorical variables.

```
gapminder %>%
  Hmisc::describe()
```

```
## .
##
## 6 Variables      1704 Observations
## -----
## country
##      n missing distinct
##    1704      0      142
##
## lowest : Afghanistan      Albania      Algeria      Angola      Argentina
## highest: Vietnam          West Bank and Gaza Yemen, Rep.      Zambia      Zimbabwe
## -----
## continent
##      n missing distinct
##    1704      0      5
##
## Value      Africa Americas      Asia      Europe      Oceania
## Frequency    624      300      396      360      24
## Proportion   0.366   0.176   0.232   0.211   0.014
## -----
## year
##      n missing distinct      Info      Mean      Gmd      .05      .10      .25      .50      .75
##    1704      0      12    0.993    1980    19.87    1952    1957    1966    1980    1993
##
## Value      1952.00 1956.95 1961.90 1966.85 1971.80 1976.75 1981.70 1986.65 1991.60 1996.55 2001.50 2007
## Frequency    142    142    142    142    142    142    142    142    142    142    142
## Proportion   0.083   0.083   0.083   0.083   0.083   0.083   0.083   0.083   0.083   0.083   0.083
##
## For the frequency table, variable is rounded to the nearest 0.55
## -----
## lifeExp
##      n missing distinct      Info      Mean      Gmd      .05      .10      .25      .50      .75
##    1704      0    1626      1    59.47    14.82    38.49    41.51    48.20    60.71    70.85
```

```
##
## lowest : 23.599 28.801 30      30.015 30.331, highest: 81.701 81.757 82      82.208 82.603
## -----
## pop
##      n missing distinct      Info      Mean      Gmd      .05      .10      .25      .50      .75
##    1704      0      1704      1 29601212 46384459  475459  946367  2793664  7023596 19585222
##
## lowest :      60011      61325      63149      65345      70787, highest: 1110396331 1164970000 123
## -----
## gdpPercap
##      n missing distinct      Info      Mean      Gmd      .05      .10      .25      .50      .75
##    1704      0      1704      1    7215    8573    548.0    687.7    1202.1    3531.8    9325.5
##
## lowest : 241.166 277.552 298.846 299.85  312.188, highest: 80894.9 95458.1 108382  109348  113523
## -----
```


Chapter 4

Introduction to Data Wrangling

In this chapter we present some very basic data handling and processing functions (data wrangling) that will be necessary for doing basic analyses, comparisons, and graphics. Most of the commands presented in this section stress the functions and R packages in the *tidyverse* - a set or family of packages that have similar syntax and behaviors.

4.1 Tidy Data

What is tidy data? Tidy data is a standard way of mapping the meaning of a dataset to its structure. A dataset is messy or tidy depending on how rows, columns and tables are matched up with observations, variables and types. In tidy data:

- Each variable forms a column.
- Each observation forms a row.
- Each type of observational unit forms a table.

4.2 Subset Observations using filter

The objective here is to examine a subset of observations, or rows of the dataset.

Suppose we wish to examine a subset of the gapminder data for only one country, Jon's favorite country, Australia!! The following code starts by taking the gapminder dataset and then “pipes” it into the filtering (selecting rows) action so that only dataset rows from Australia are selected. The pipe function is `%>%` and is similar to a plumbing pipe that goes one direction: from left to right. After the “Australia” rows are selected, the result is “piped” into the head function for display. The head function says show the top 12 rows. When no rows are specified in the head function, the default is 6 rows. Note that the filter function resides in the *dplyr* package within the *tidyverse* family.

If the *tidyverse* or *dplyr* packages have been loaded, you don't need to supply the `dplyr::` prefix to the `filter` command.

```
gapminder %>% dplyr::filter(country=="Australia") %>% head(n=12)
```

```
## # A tibble: 12 x 6
##   country    continent  year lifeExp      pop gdpPercap
##   <fct>      <fct>    <int>  <dbl>    <int>    <dbl>
## 1 Australia Oceania    1952   69.1  8691212  10040.
## 2 Australia Oceania    1957   70.3  9712569  10950.
```

```
## 3 Australia Oceania 1962 70.9 10794968 12217.
## 4 Australia Oceania 1967 71.1 11872264 14526.
## 5 Australia Oceania 1972 71.9 13177000 16789.
## 6 Australia Oceania 1977 73.5 14074100 18334.
## 7 Australia Oceania 1982 74.7 15184200 19477.
## 8 Australia Oceania 1987 76.3 16257249 21889.
## 9 Australia Oceania 1992 77.6 17481977 23425.
## 10 Australia Oceania 1997 78.8 18565243 26998.
## 11 Australia Oceania 2002 80.4 19546792 30688.
## 12 Australia Oceania 2007 81.2 20434176 34435.
```

```
# will produce the same output as
#gapminder %>% filter(country=="Australia") %>% head(n=12)
```

4.3 Subset using multiple conditions

Let's select observations by continent and year. The `head` function will then show some of the rows selected. Here the `gapminder` dataframe is piped to the `filter` function to select rows to be further piped to the `head()` function for display. The logical condition inside `filter` restricts continent to "Oceania" AND (AND condition is "&") year to be 1997. Both of these conditions must be TRUE for the row to enter the dataframe to displayed by the `head()` function.

```
gapminder %>%
  dplyr::filter(continent=="Oceania" & year==1997) %>%
  head()

## # A tibble: 2 x 6
##   country    continent  year lifeExp      pop gdpPercap
##   <fct>      <fct>    <int>  <dbl>    <int>    <dbl>
## 1 Australia  Oceania   1997   78.8 18565243  26998.
## 2 New Zealand Oceania   1997   77.6  3676187  21050.
```

Notice that two filter statements produce the same result.

```
gapminder %>%
  dplyr::filter(continent=="Oceania") %>%
  dplyr::filter(year==1997) %>%
  head()

## # A tibble: 2 x 6
##   country    continent  year lifeExp      pop gdpPercap
##   <fct>      <fct>    <int>  <dbl>    <int>    <dbl>
## 1 Australia  Oceania   1997   78.8 18565243  26998.
## 2 New Zealand Oceania   1997   77.6  3676187  21050.
```

The next example uses an "or" condition to specify the desired rows in the first `filter` expression - the next `filter` permits only observations from 1997.

```
gapminder %>%
  dplyr::filter(continent=="Oceania" | continent=="Americas") %>%
  dplyr::filter(year==1997) %>%
  head()

## # A tibble: 6 x 6
##   country    continent  year lifeExp      pop gdpPercap
##   <fct>      <fct>    <int>  <dbl>    <int>    <dbl>
## 1 Argentina  Americas  1997   73.3 36203463  10967.
## 2 Australia  Oceania   1997   78.8 18565243  26998.
```

```
## 3 Bolivia Americas 1997 62.0 7693188 3326.
## 4 Brazil Americas 1997 69.4 168546719 7958.
## 5 Canada Americas 1997 78.6 30305843 28955.
## 6 Chile Americas 1997 75.8 14599929 10118.
```

The next example selects observations/rows from a list of countries and also restricts year to 1997.

```
gapminder %>%
  filter(country %in% c("Australia", "New Zealand", "Argentina") & year==1997) %>%
  head()
```

```
## # A tibble: 3 x 6
##   country    continent  year lifeExp      pop gdpPercap
##   <fct>      <fct>    <int> <dbl>    <int>    <dbl>
## 1 Argentina Americas   1997  73.3 36203463  10967.
## 2 Australia Oceania    1997  78.8 18565243  26998.
## 3 New Zealand Oceania    1997  77.6  3676187  21050.
```

The next example selects observations by omitting one continent (Oceania is excluded) and then specifies a year. The code that causes “omit” is the “!=” syntax. In the code `year==1997`, the double equal sign `==` means make a logical check if year is 1997. Only rows where both aspects of the filter conditions pass through to be displayed by `head`. Again, the logical operator “AND” is expressed by the `&` expression.

```
gapminder %>%
  filter(continent!="Oceania" & year==1997) %>%
  head()
```

```
## # A tibble: 6 x 6
##   country    continent  year lifeExp      pop gdpPercap
##   <fct>      <fct>    <int> <dbl>    <int>    <dbl>
## 1 Afghanistan Asia      1997  41.8 22227415    635.
## 2 Albania Europe    1997  73.0  3428038   3193.
## 3 Algeria Africa    1997  69.2 29072015   4797.
## 4 Angola Africa    1997  41.0  9875024   2277.
## 5 Argentina Americas   1997  73.3 36203463  10967.
## 6 Austria Europe    1997  77.5  8069876  29096.
```

Please note that in all the above examples, the `filter` function accepts/rejects rows or observations in a dataframe according to the logical conditions specified inside the filter function.

4.4 Saving as a new dataframe

Here we save the the modified dataset as a new dataframe called `gap97` by using an assignment arrow “`<-`” to save the resulting pipeline to `gap97`.

```
gap97 <- gapminder %>%
  filter(continent!="Oceania" & year==1997)
#
dplyr::glimpse(gap97)
```

```
## Rows: 140
## Columns: 6
## $ country    <fct> "Afghanistan", "Albania", "Algeria", "Angola", "Argentina", "Austria", "Bahrain", "Ba
## $ continent  <fct> Asia, Europe, Africa, Africa, Americas, Europe, Asia, Asia, Europe, Africa, Americas,
## $ year       <int> 1997, 1997, 1997, 1997, 1997, 1997, 1997, 1997, 1997, 1997, 1997, 1997, 1997, 1997, 1997, 1
## $ lifeExp    <dbl> 41.763, 72.950, 69.152, 40.963, 73.275, 77.510, 73.925, 59.412, 77.530, 54.777, 62.05
## $ pop        <int> 22227415, 3428038, 29072015, 9875024, 36203463, 8069876, 598561, 123315288, 10199787,
## $ gdpPercap  <dbl> 635.3414, 3193.0546, 4797.2951, 2277.1409, 10967.2820, 29095.9207, 20292.0168, 972.77
```

4.5 Subset using top_n

Let's make a dataset based on the ten countries in 1997 with highest gdp.

```
gapminder %>%
  filter(year==1997) %>%
  top_n(n = 10, wt = gdpPercap) %>%
  head(n=10)
```

```
## # A tibble: 10 x 6
##   country      continent  year lifeExp      pop gdpPercap
##   <fct>        <fct>    <int>  <dbl>    <int>    <dbl>
## 1 Austria      Europe    1997   77.5  8069876  29096.
## 2 Canada      Americas  1997   78.6  30305843 28955.
## 3 Denmark      Europe    1997   76.1  5283663  29804.
## 4 Japan        Asia     1997   80.7 125956499 28817.
## 5 Kuwait       Asia     1997   76.2  1765345  40301.
## 6 Netherlands  Europe    1997   78.0  15604464  30246.
## 7 Norway       Europe    1997   78.3  4405672  41283.
## 8 Singapore    Asia     1997   77.2  3802309  33519.
## 9 Switzerland  Europe    1997   79.4  7193761  32135.
## 10 United States Americas  1997   76.8 272911760 35767.
```

4.6 Subset using select

The `filter` function relates to the rows of the dataframe. Sometimes we might want to include only a few of the variables (columns) in a dataset. We frequently want to create a data subset with only a few variables when the original dataset has many variables. The `select` function is used to select and rename variables.

the next command selects three variables and renames two of them:

```
gapminder %>%
  dplyr::select(country, Year=year, LifeExp=lifeExp) %>%
  head()
```

```
## # A tibble: 6 x 3
##   country      Year LifeExp
##   <fct>        <int>  <dbl>
## 1 Afghanistan  1952   28.8
## 2 Afghanistan  1957   30.3
## 3 Afghanistan  1962   32.0
## 4 Afghanistan  1967   34.0
## 5 Afghanistan  1972   36.1
## 6 Afghanistan  1977   38.4
```

to change the order of display, puts year first in the list of variables

```
gapminder %>%
  select(year, everything()) %>%
  head()
```

```
## # A tibble: 6 x 6
##   year country      continent lifeExp      pop gdpPercap
##   <int> <fct>        <fct>    <dbl>    <int>    <dbl>
## 1  1952 Afghanistan Asia      28.8  8425333    779.
## 2  1957 Afghanistan Asia      30.3  9240934    821.
## 3  1962 Afghanistan Asia      32.0 10267083    853.
## 4  1967 Afghanistan Asia      34.0 11537966    836.
```

```
## 5 1972 Afghanistan Asia      36.1 13079460      740.
## 6 1977 Afghanistan Asia      38.4 14880372      786.
```

The `profiling_num` command from the *funModeling* package produces a lot of output, some we might not want. We will show how to modify the output of this command here. The command produces a dataframe which has many columns we might not wish to display or consider further.

We begin by removing some columns of summary statistics that we wish to ignore. Selecting a list of column names with a “minus” - sign in front of the list will remove these items from the dataframe and keep the rest in place. The command below pipes the modified dataframe to the `kable` command in the *knitr* package for a more pleasing tabular display.

```
# Let's observe the contents of profiling_num:
```

```
funModeling::profiling_num(gapminder) %>%
  dplyr::glimpse()
```

```
## Rows: 4
## Columns: 16
## $ variable      <chr> "year", "lifeExp", "pop", "gdpPercap"
## $ mean          <dbl> 1.979500e+03, 5.947444e+01, 2.960121e+07, 7.215327e+03
## $ std_dev       <dbl> 1.726533e+01, 1.291711e+01, 1.061579e+08, 9.857455e+03
## $ variation_coef <dbl> 0.008722066, 0.217187544, 3.586268548, 1.366182632
## $ p_01          <dbl> 1952.0000, 33.4926, 154117.9200, 369.2201
## $ p_05          <dbl> 1952.0000, 38.4924, 475458.9000, 547.9964
## $ p_25          <dbl> 1965.750, 48.198, 2793664.000, 1202.060
## $ p_50          <dbl> 1979.5000, 60.7125, 7023595.5000, 3531.8470
## $ p_75          <dbl> 1.993250e+03, 7.084550e+01, 1.958522e+07, 9.325462e+03
## $ p_95          <dbl> 2007.000, 77.437, 89822054.500, 26608.333
## $ p_99          <dbl> 2.007000e+03, 8.023892e+01, 6.319900e+08, 3.678357e+04
## $ skewness      <dbl> 0.0000000, -0.2524798, 8.3328742, 3.8468819
## $ kurtosis      <dbl> 1.783217, 1.873099, 80.716151, 30.431702
## $ iqr           <dbl> 2.750000e+01, 2.264750e+01, 1.679156e+07, 8.123402e+03
## $ range_98      <chr> "[1952, 2007]", "[33.4926, 80.23892]", "[154117.92, 631990000.000002]", "[369.2201, 16791557.75]"
## $ range_80      <chr> "[1957, 2002]", "[41.5108, 75.097]", "[946367.1, 54801369.5]", "[687.71836128, 16791557.75]"
```

```
# now remove unwanted columns from summary display
```

```
funModeling::profiling_num(gapminder) %>%
  select(-c("variation_coef", "skewness", "kurtosis", "range_98", "range_80", "p_01", "p_99")) %>%
  knitr::kable(., digits=2)
```

variable	mean	std_dev	p_05	p_25	p_50	p_75	p_95	iqr
year	1979.50	17.27	1952.00	1965.75	1979.50	1993.25	2007.00	27.50
lifeExp	59.47	12.92	38.49	48.20	60.71	70.85	77.44	22.65
pop	29601212.32	106157896.74	475458.90	2793664.00	7023595.50	19585221.75	89822054.50	16791557.75
gdpPercap	7215.33	9857.45	548.00	1202.06	3531.85	9325.46	26608.33	8123.40

In the next command we take a different approach - we explicitly select the statistics (columns) we want to keep and display. The most commonly used summaries are chosen.

```
funModeling::profiling_num(gapminder) %>%
  select(c("variable", "mean", "std_dev", "p_25", "p_50", "p_75")) %>%
  knitr::kable(., digits=2)
```

variable	mean	std_dev	p_25	p_50	p_75
year	1979.50	17.27	1965.75	1979.50	1993.25
lifeExp	59.47	12.92	48.20	60.71	70.85
pop	29601212.32	106157896.74	2793664.00	7023595.50	19585221.75
gdpPercap	7215.33	9857.45	1202.06	3531.85	9325.46

4.7 Order using arrange

Sometimes we might want to know the countries with the largest or smallest values of some variables. In the following examples we sort/order by the values of life expectancy. In the code below, when we use the command `filter(year==1997)`, the double equal sign means make a logical check if year is 1997, and only allow dataframe rows where this is true to pass through to the next stage of the analysis pipeline. The `desc` function means the values will be arranged in descending values (large to small) - default is ascending (small to large).

```
# This command will show the countries with highest life expectancy because
# the data are arranged in descending order of life expectancy (larger to smaller)
gapminder %>%
  dplyr::filter(year==1997) %>%
  dplyr::select(country, continent, lifeExp) %>%
  dplyr::arrange(desc(lifeExp)) %>%
  head()
```

```
## # A tibble: 6 x 3
##   country      continent lifeExp
##   <fct>        <fct>      <dbl>
## 1 Japan        Asia          80.7
## 2 Hong Kong, China Asia          80
## 3 Sweden       Europe         79.4
## 4 Switzerland Europe         79.4
## 5 Iceland      Europe         79.0
## 6 Australia    Oceania        78.8
```

```
# This command uses the default ascending (increasing) order with
# respect to life expectancy (order smaller to larger)
gapminder %>%
  filter(year==1997) %>%
  select(country, continent, lifeExp) %>%
  arrange(lifeExp) %>%
  head()
```

```
## # A tibble: 6 x 3
##   country      continent lifeExp
##   <fct>        <fct>      <dbl>
## 1 Rwanda      Africa         36.1
## 2 Sierra Leone Africa         39.9
## 3 Zambia      Africa         40.2
## 4 Angola      Africa         41.0
## 5 Afghanistan Asia          41.8
## 6 Liberia     Africa         42.2
```

The `top_n` function from the *dplyr* package will select the `n` rows with the largest values of a variable. This is similar to the code above that orders the rows - then use `head` function to select the number of desired rows.

This first example uses the default alphabetical ordering of country name.

```
gapminder %>%
  filter(year==1997) %>%
  select(country, continent, lifeExp) %>%
  dplyr::top_n(n=6, wt=lifeExp) %>%
  knitr::kable(., digits=2)
```

country	continent	lifeExp
Australia	Oceania	78.83
Hong Kong, China	Asia	80.00
Iceland	Europe	78.95
Japan	Asia	80.69
Sweden	Europe	79.39
Switzerland	Europe	79.37

The results can then be ordered by the life expectancy:

```
gapminder %>%
  filter(year==1997) %>%
  select(country, continent, lifeExp) %>%
  dplyr::top_n(n=6, wt=lifeExp) %>%
  dplyr::arrange(desc(lifeExp)) %>%
  knitr::kable(., digits=2)
```

country	continent	lifeExp
Japan	Asia	80.69
Hong Kong, China	Asia	80.00
Sweden	Europe	79.39
Switzerland	Europe	79.37
Iceland	Europe	78.95
Australia	Oceania	78.83

The countries with the largest life expectancy can then be ordered by another variable like population. Here we find the 6 countries in 1997 with the highest life expectancy - then display them in order of population size.

```
gapminder %>%
  filter(year==1997) %>%
  select(country, continent, lifeExp, pop) %>%
  dplyr::top_n(n=6, wt=lifeExp) %>%
  dplyr::arrange(desc(pop)) %>%
  knitr::kable(., digits=2)
```

country	continent	lifeExp	pop
Japan	Asia	80.69	125956499
Australia	Oceania	78.83	18565243
Sweden	Europe	79.39	8897619
Switzerland	Europe	79.37	7193761
Hong Kong, China	Asia	80.00	6495918
Iceland	Europe	78.95	271192

4.8 Grouped Filter

Another useful verb in the *tidyverse* is `group_by`. Suppose we wanted to view the two countries with the highest life expectancy in 1997, in each continent.

```
gapminder %>%
  filter(year==1997) %>%
  select(country, continent, lifeExp, pop) %>%
  dplyr::group_by(continent) %>%
  dplyr::top_n(n=2, wt=lifeExp) %>%
  dplyr::arrange(continent) %>%
  knitr::kable()
```

country	continent	lifeExp	pop
Reunion	Africa	74.772	684810
Tunisia	Africa	71.973	9231669
Canada	Americas	78.610	30305843
Costa Rica	Americas	77.260	3518107
Hong Kong, China	Asia	80.000	6495918
Japan	Asia	80.690	125956499
Sweden	Europe	79.390	8897619
Switzerland	Europe	79.370	7193761
Australia	Oceania	78.830	18565243
New Zealand	Oceania	77.550	3676187

4.9 New Variables Using mutate

In many problems we may wish to create a new variable based on an existing variable. Here we illustrate by making a new variable - the natural logarithm of population - based on the original variable `pop` using the `mutate` function.

```
gapminder %>%
  dplyr::mutate(logpopulation = log(pop)) %>%
  dplyr::glimpse()

## Rows: 1,704
## Columns: 7
## $ country      <fct> "Afghanistan", "Afghanistan", "Afghanistan", "Afghanistan", "Afghanistan", "A
## $ continent    <fct> Asia, Asia, Asia, Asia, Asia, Asia, Asia, Asia, Asia, Asia, Asia, Asia, Europ
## $ year         <int> 1952, 1957, 1962, 1967, 1972, 1977, 1982, 1987, 1992, 1997, 2002, 2007, 1952,
## $ lifeExp      <dbl> 28.801, 30.332, 31.997, 34.020, 36.088, 38.438, 39.854, 40.822, 41.674, 41.76
## $ pop          <int> 8425333, 9240934, 10267083, 11537966, 13079460, 14880372, 12881816, 13867957,
## $ gdpPercap    <dbl> 779.4453, 820.8530, 853.1007, 836.1971, 739.9811, 786.1134, 978.0114, 852.395
## $ logpopulation <dbl> 15.94675, 16.03915, 16.14445, 16.26115, 16.38655, 16.51555, 16.37133, 16.4450
```

If I want to change the name of the new variable from `logpopulation` to something shorter like `logPop`, we could re-run the `mutate` command, or use a `rename` function.

In addition we create a new version of the `gapminder` dataset that contains the new variable - called `gapVers1`. This dataframe is now available to be used in the ongoing analysis.

```
gapVers1 <- gapminder %>%
  dplyr::mutate(logpopulation = log(pop)) %>%
  dplyr::rename(logPop=logpopulation)
#
dplyr::glimpse(gapVers1)

## Rows: 1,704
## Columns: 7
## $ country      <fct> "Afghanistan", "Afghanistan", "Afghanistan", "Afghanistan", "Afghanistan", "Afgha
## $ continent    <fct> Asia, Asia, Asia, Asia, Asia, Asia, Asia, Asia, Asia, Asia, Asia, Asia, Europe, E
## $ year         <int> 1952, 1957, 1962, 1967, 1972, 1977, 1982, 1987, 1992, 1997, 2002, 2007, 1952, 195
## $ lifeExp      <dbl> 28.801, 30.332, 31.997, 34.020, 36.088, 38.438, 39.854, 40.822, 41.674, 41.763, 4
## $ pop          <int> 8425333, 9240934, 10267083, 11537966, 13079460, 14880372, 12881816, 13867957, 163
## $ gdpPercap    <dbl> 779.4453, 820.8530, 853.1007, 836.1971, 739.9811, 786.1134, 978.0114, 852.3959, 6
## $ logPop       <dbl> 15.94675, 16.03915, 16.14445, 16.26115, 16.38655, 16.51555, 16.37133, 16.44509, 1
```

The next code uses a `mutate` command with logical conditions to make a new, two-level categorical variable `region` as a character variable. Then we use `mutate` again to convert `region` (character) to a factor variable


```
dplyr::filter(continent=="Americas") %>%
dplyr::count()
```

```
## # A tibble: 1 x 1
##       n
##   <int>
## 1     25
```

Now we group by continent.

```
gapminder %>%
  dplyr::filter(year==1997) %>%
  dplyr::group_by(continent) %>%
  dplyr::filter(continent=="Americas") %>%
  dplyr::tally()
```

```
## # A tibble: 1 x 2
##   continent     n
##   <fct>       <int>
## 1 Americas     25
```

```
#
gapminder %>%
  dplyr::filter(year==1997) %>%
  dplyr::group_by(continent) %>%
  dplyr::tally()
```

```
## # A tibble: 5 x 2
##   continent     n
##   <fct>       <int>
## 1 Africa      52
## 2 Americas    25
## 3 Asia       33
## 4 Europe     30
## 5 Oceania     2
```

```
gapminder %>%
  dplyr::filter(year==1997) %>%
  dplyr::group_by(continent) %>%
  dplyr::filter(continent=="Americas") %>%
  dplyr::count()
```

```
## # A tibble: 1 x 2
## # Groups:   continent [1]
##   continent     n
##   <fct>       <int>
## 1 Americas     25
```

```
#
gapminder %>%
  dplyr::filter(year==1997) %>%
  dplyr::count(continent)
```

```
## # A tibble: 5 x 2
##   continent     n
##   <fct>       <int>
## 1 Africa      52
## 2 Americas    25
## 3 Asia       33
```

```
## 4 Europe      30
## 5 Oceania     2
```

4.11 Missing Values

If a variable is not complete and contains empty places, these are denoted in R as NA. We will often wish to create a dataframe without any missing values, or discover how many rows contain variables with missing values.

First let's create a small dataset with missing values:

```
x <- c(1,2,NA,4)
y <- c(11,12,13,NA)
z <- c(7,8,9,10)
tempdf <- data.frame(x,y,z)
tempdf
```

```
##      x  y  z
## 1   1 11  7
## 2   2 12  8
## 3  NA 13  9
## 4   4 NA 10
```

```
# count missing values for variable x
tempdf %>%
  dplyr::summarise(count = sum(is.na(x)))
```

```
##      count
## 1         1
```

```
# count rows with missing values for variable y
tempdf %>%
  dplyr::tally(is.na(y))
```

```
##      n
## 1 1
```

```
# subset of rows with complete data for specified columns
tempdf %>%
  dplyr::select(y,z) %>%
  tidyr::drop_na() %>%
  head()
```

```
##      y  z
## 1 11  7
## 2 12  8
## 3 13  9
```

```
# drop rows with missing values in all variables
tempdf %>%
  tidyr::drop_na() %>%
  head()
```

```
##      x  y  z
## 1 1 11  7
## 2 2 12  8
```

Use base `is.na` function

```
tempdf %>%  
  filter(!is.na(x), # allow obs with nonmissing x  
         !is.na(y), # allow obs with nonmissing y  
         !is.na(z)) # allow obs with nonmissing z
```

```
##    x  y z  
##  1 1 11 7  
##  2 2 12 8
```

Next, some code that will execute a filter that will permit only rows with entirely complete data in x to pass through to the dataset,

```
tempdf %>%  
  filter(x %>% is.na() %>% magrittr::not()) %>%  
  head()
```

```
##    x  y z  
##  1 1 11 7  
##  2 2 12 8  
##  3 4 NA 10
```

Chapter 5

Univariate Graphical Displays

In this section we will show examples of how to create graphical displays of a single variable - with examples for both quantitative and categorical variables. In each example, the first line creates the dataset to be graphed - followed by a command making the display. We will focus on graphical displays made by functions in the *ggplot2* family - that is, the *ggplot2* package which is also part of the *tidyverse* family of functions. If *tidyverse* is loaded, *ggplot2* functions will work without explicitly loading the *ggplot2* package.

5.1 Overview of ggplot

The *ggplot2* package uses the **ggplot** command - and builds a graphical display in steps and layers. We always start with the **ggplot** command which typically has two basic elements: a dataset to be used, and a list of mappings using the **aes** instruction used to connect dataset variables to aspects of the plot like the vertical axis, horizontal axis, or perhaps the size of a point.


The kind of object being displayed is called a **geom**, (a geometry) and a plot can have several **geoms**, and they are added to a display in layers - connected by a **+** sign.

5.2 A Quantitative Variable

5.2.1 Dotplot

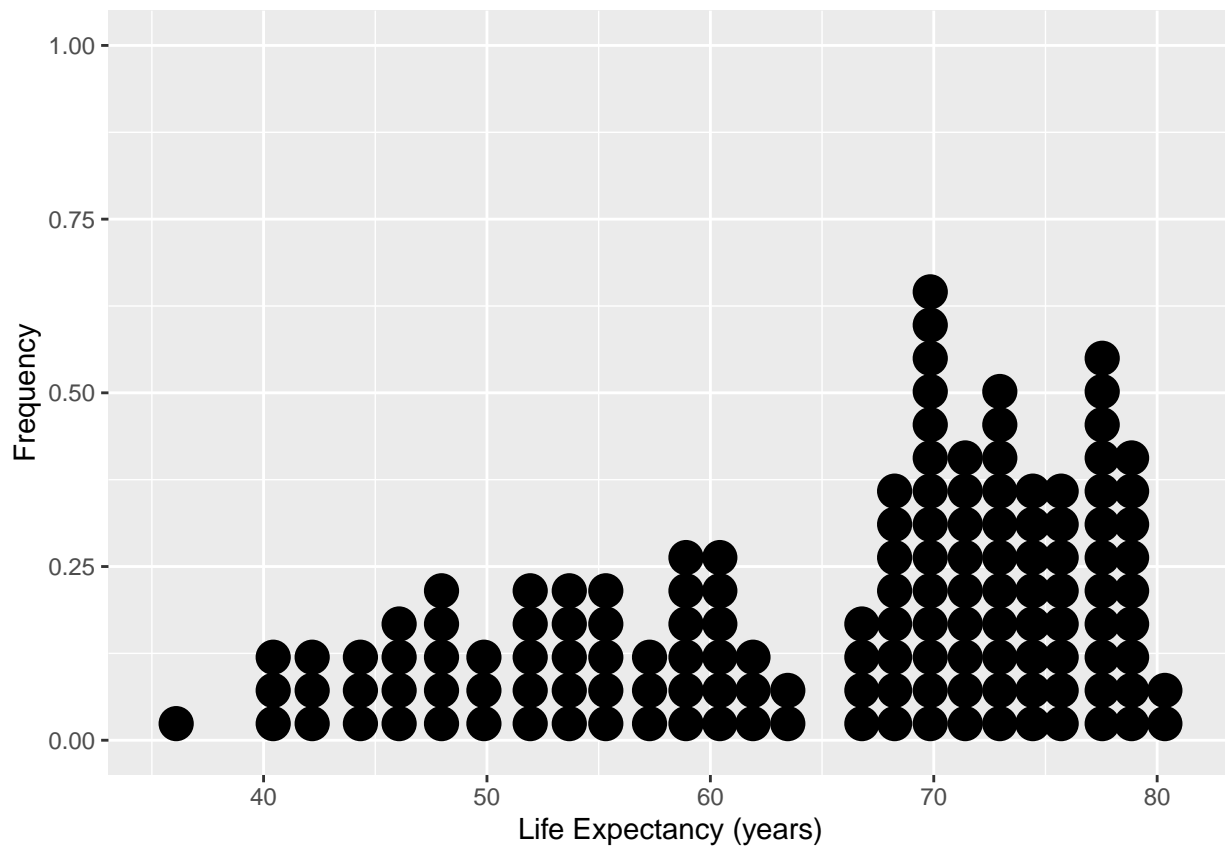
The next block of code takes the *gapminder* dataframe and “pipes” (**%>%**, a pipeline like plumbing) the data through a filter so that only data from year 1997 flows through to define the new dataset named **ds**. The **ggplot** command uses dataset **ds**, and variable **x** life expectancy. The next example shows what using only the **ggplot** command produces an empty graphical region that is awaiting further instructions:

```
ds <- gapminder %>%  
  filter(year==1997)  
#  
ggplot(data=ds, mapping=aes(x=lifeExp))
```



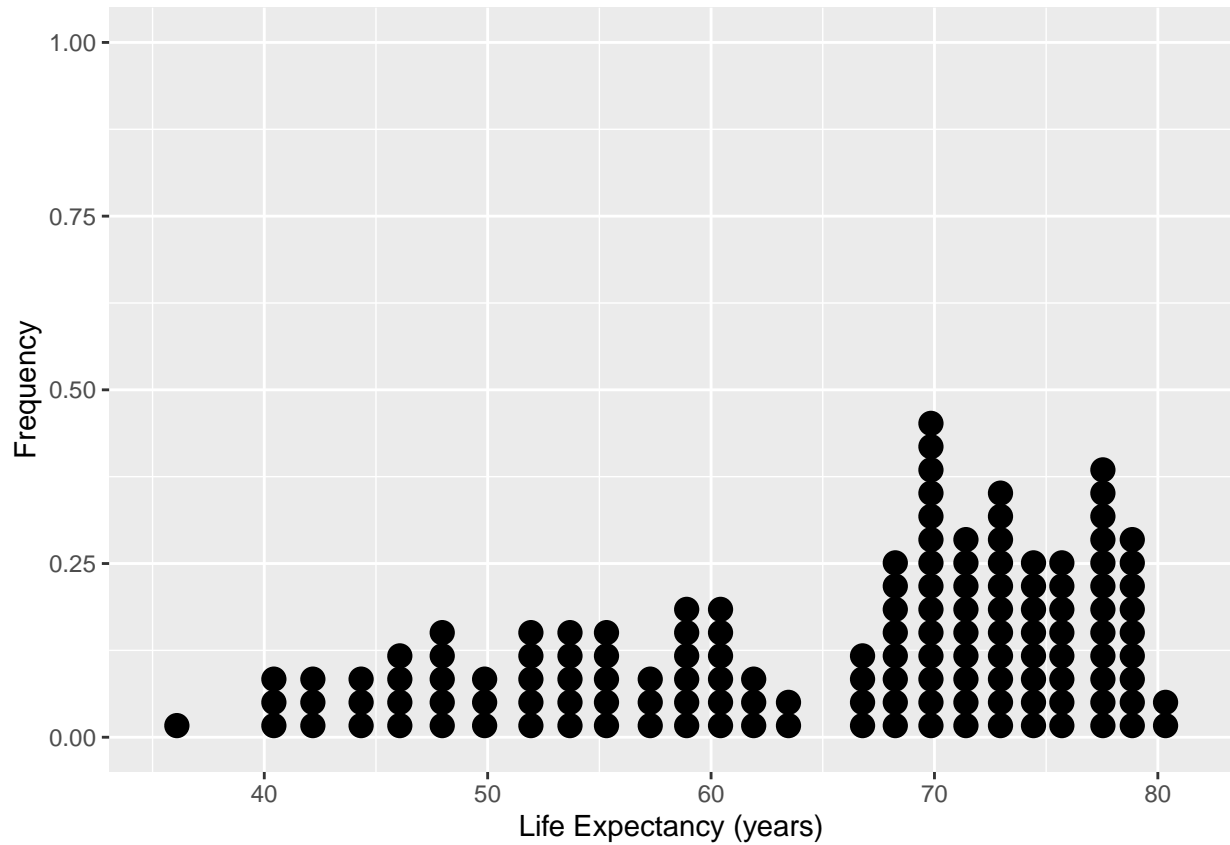
Now we use additional code to place the dotplot in the existing graphical region. In ggplot graphics we make graphical objects with a `geom` function - here a dotplot so we use `geom_dotplot()` to produce the dotplot specified using the variable mappings in the aesthetics command `aes` in the `ggplot` command.

```
ds <- gapminder %>%  
  filter(year==1997)  
#  
ggplot(data=ds, mapping=aes(x=lifeExp)) +  
  geom_dotplot() +  
  xlab("Life Expectancy (years)") + ylab("Frequency")
```



Here we change the default size for the dots, and pipe the data directly into the first argument data (using `data=.`) of the `ggplot` command:

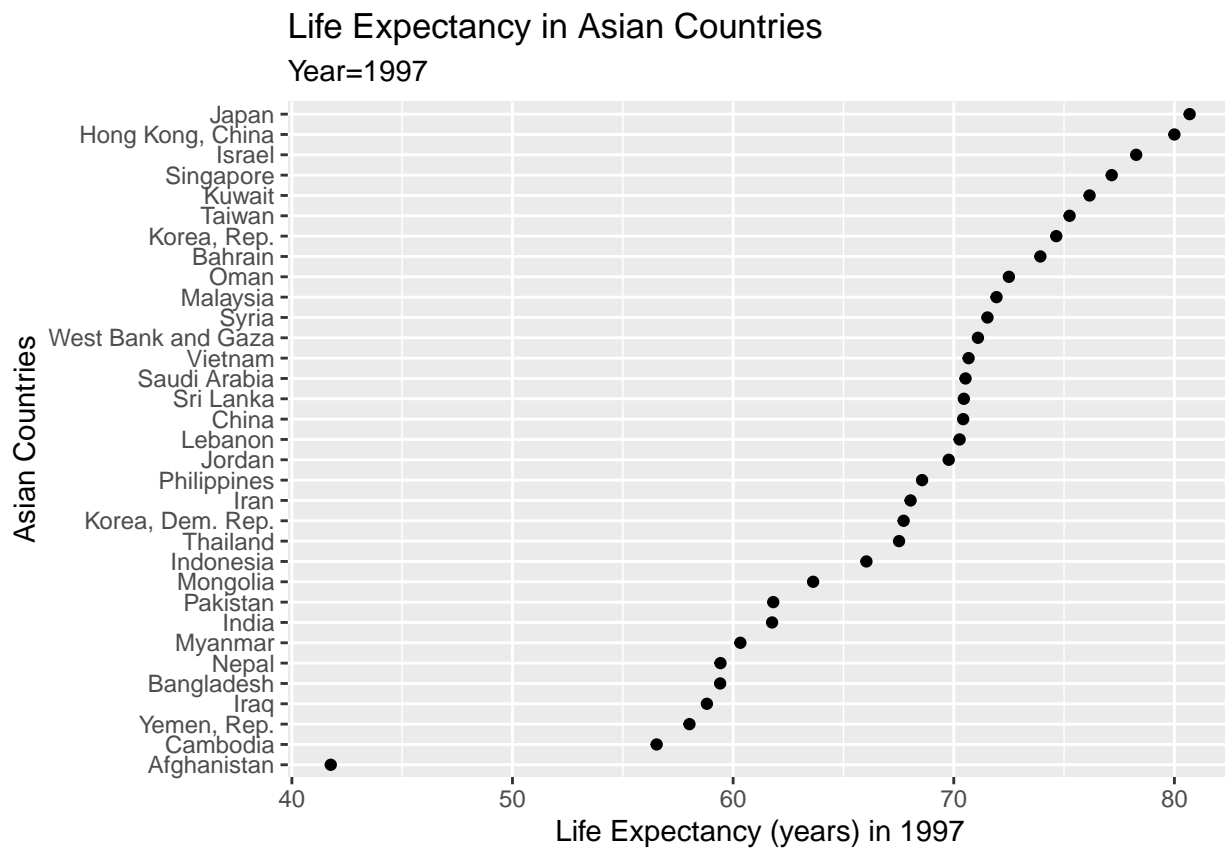
```
gapminder %>%
  filter(year==1997) %>%
  ggplot(data=., mapping=aes(x=lifeExp)) +
  geom_dotplot(dotsize=0.70) +
  xlab("Life Expectancy (years)") + ylab("Frequency")
```



5.2.1.1 Dotplot with observations identified and ordered


Here we produce a display so that life expectancy is displayed for each country in Asia, and the values are ordered.

```
ds <- gapminder %>%
  filter(continent=="Asia",year==1997)
#
ggplot(data=ds, mapping=aes(x=lifeExp, y= reorder(country,lifeExp))) +
  geom_point() +
  labs(title="Life Expectancy in Asian Countries",
       subtitle="Year=1997",
       x="Life Expectancy (years) in 1997",
       y="Asian Countries")
```

Notice that in the next example we simply pipe the modified dataset into the first argument of the `ggplot` command so that there is no need to save the modified dataset to make the display.

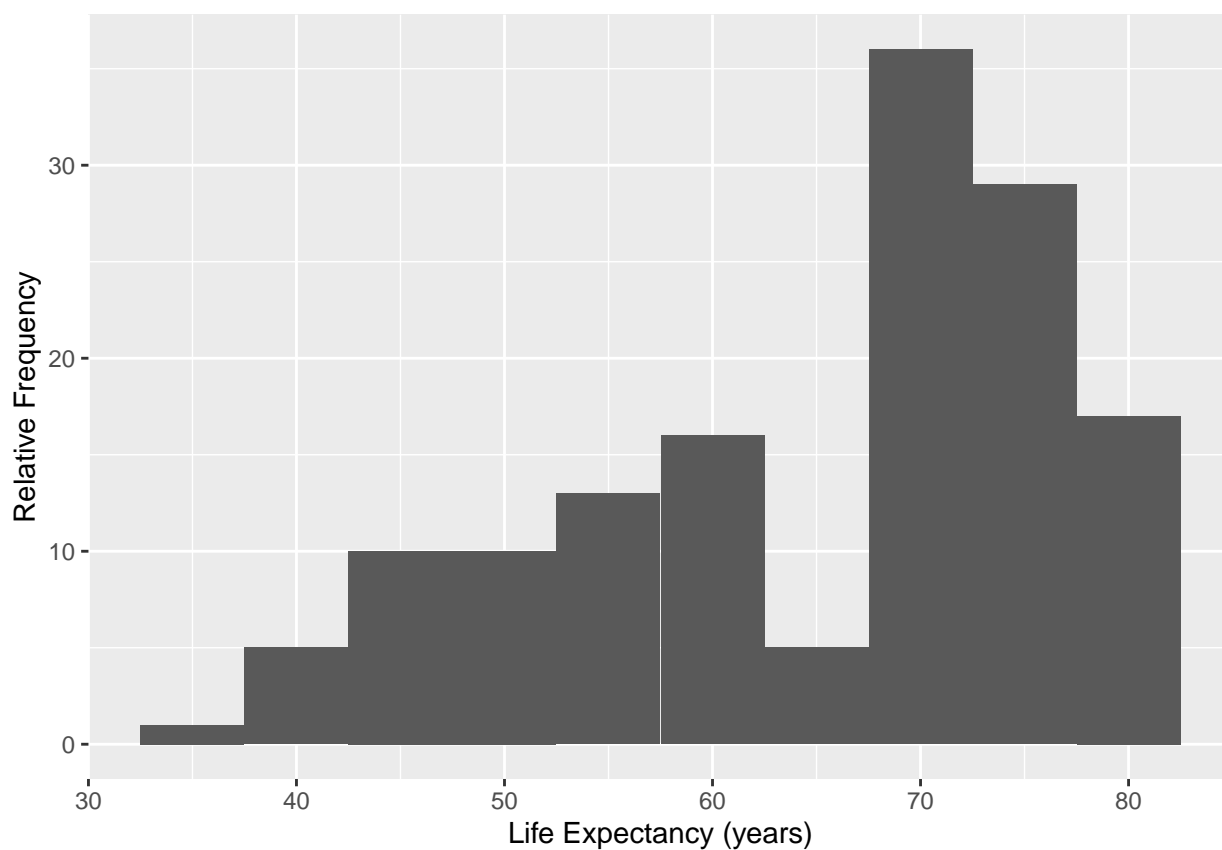
```
gapminder %>%
  filter(continent=="Asia",year==1997) %>%
  ggplot(data=., mapping=aes(x=lifeExp, y= reorder(country,lifeExp))) +
    geom_point() +
    xlab("Life Expectancy (years) in 1997") +
    ylab("Asian Countries")
```



5.2.2 Histogram

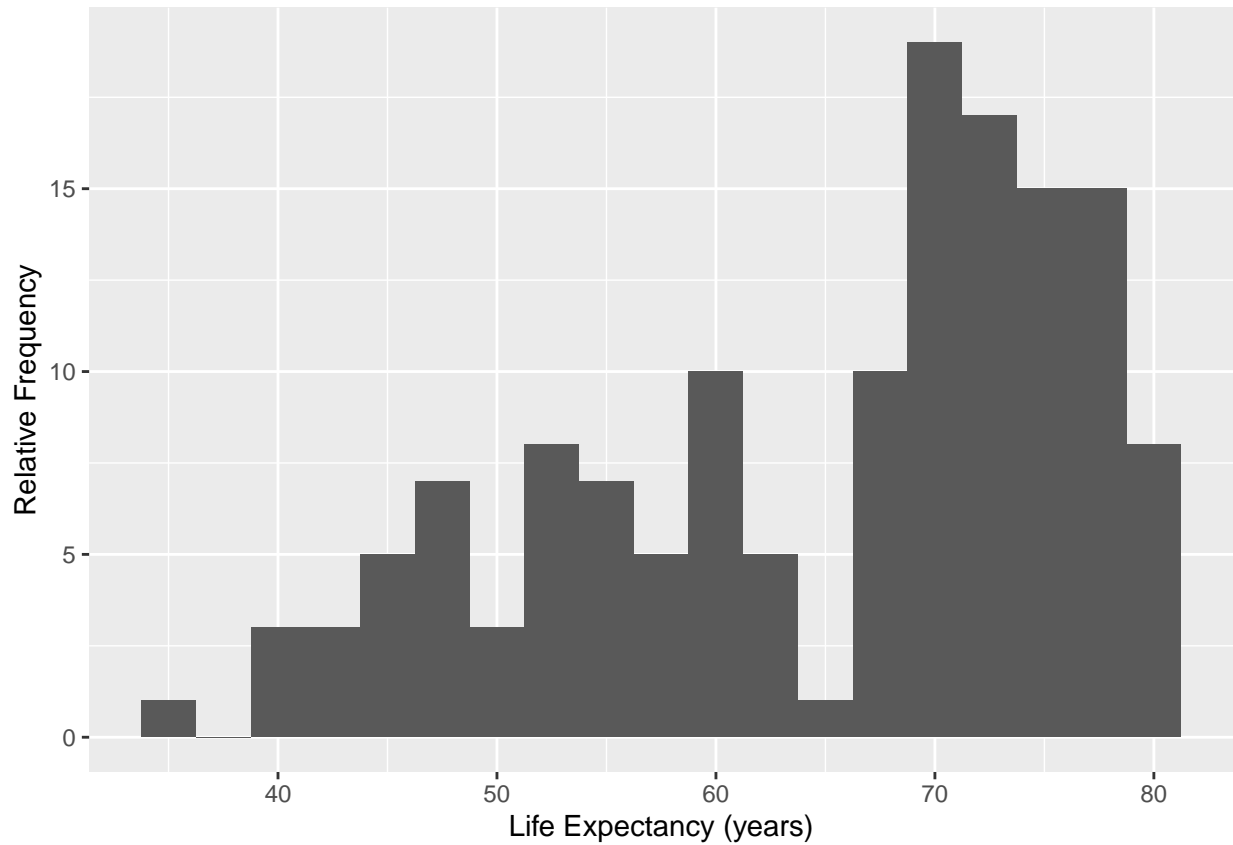
This code block is similar to the dotplot commands, but the `geom_histogram` function controls the bin width in units of the x variable - in this case 5 years.

```
gapminder %>%
  filter(year==1997) %>%
  ggplot(data=., mapping=aes(x=lifeExp)) +
    geom_histogram(binwidth=5) +
    xlab("Life Expectancy (years)") +
    ylab("Relative Frequency")
```




Here we change the binwidth:

```
gapminder %>%  
  filter(year==1997) %>%  
  ggplot(data=., mapping=aes(x=lifeExp)) +  
    geom_histogram(binwidth=2.5) +  
    xlab("Life Expectancy (years)") +  
    ylab("Relative Frequency")
```



Here we change the number of bins:

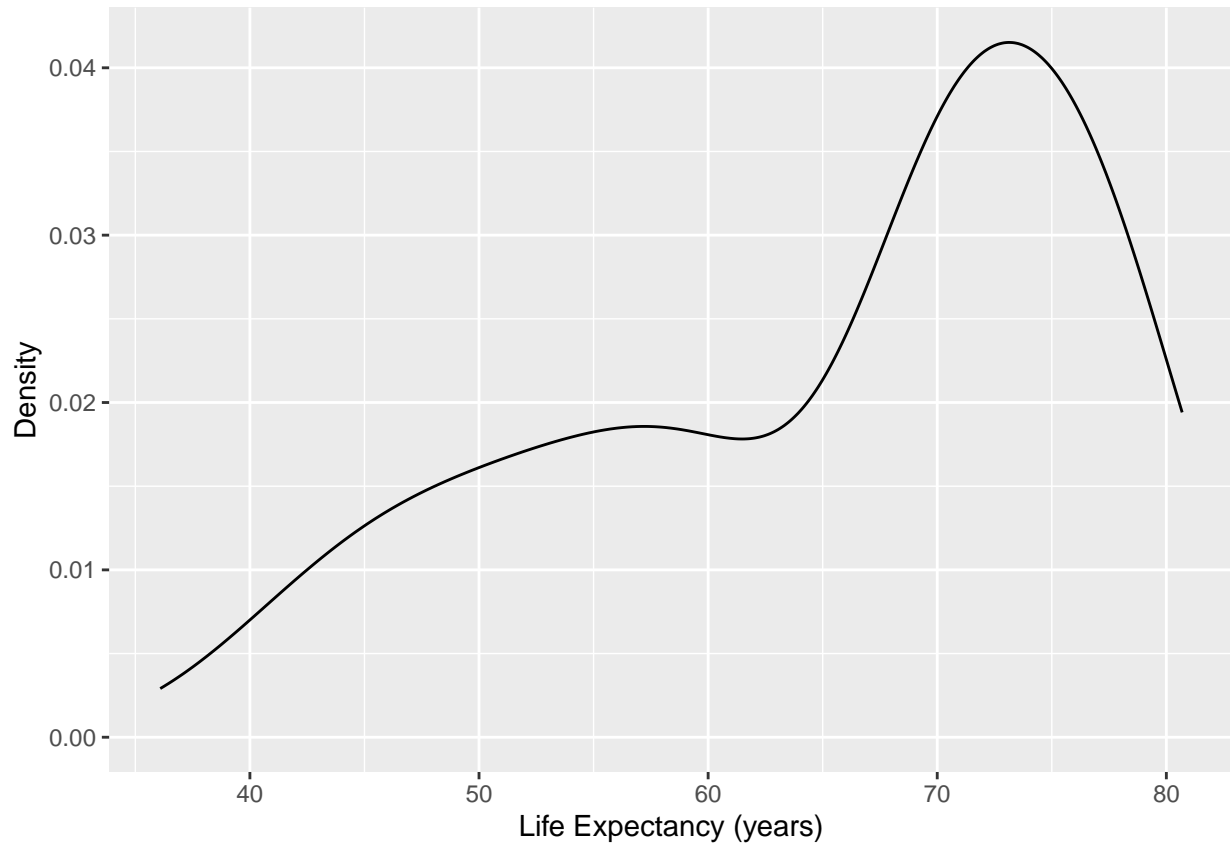
```
gapminder %>%  
  filter(year==1997) %>%  
  ggplot(data=., mapping=aes(x=lifeExp)) +  
    geom_histogram(bins=10) +  
    xlab("Life Expectancy (years)") +  
    ylab("Relative Frequency")
```



5.2.3 Density Plot


Density plots produces a smoothing of a histogram to display the distribution.

```
ds <- gapminder %>%  
  filter(year==1997)  
#  
ggplot(data=ds, mapping=aes(x=lifeExp)) +  
  geom_density() +  
  xlab("Life Expectancy (years)") +  
  ylab("Density")
```



The `adjust` option controls the amount of smoothing relative to a default value of 1. A smaller value gives less smoothing (more responsive line to small changes in the data distribution), and larger values will make a smoother curve that is less sensitive to the data pattern.

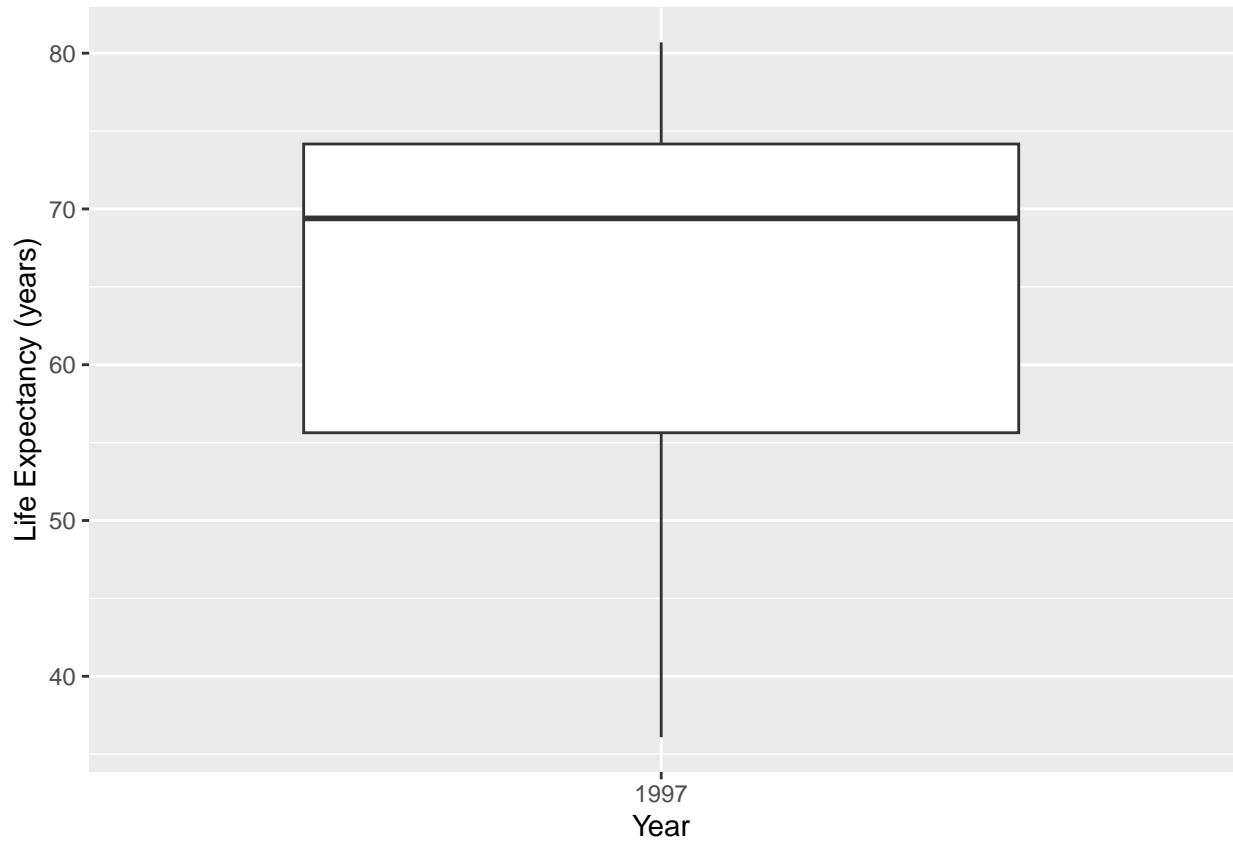
```
ds <- gapminder %>%  
  filter(year==1997)  
#  
ggplot(data=ds, mapping=aes(x=lifeExp)) +  
  geom_density(adjust=0.75) +  
  xlab("Life Expectancy (years)") +  
  ylab("Density")
```




5.2.4 Boxplot

The boxplot display really needs only a single quantitative variable (here life expectancy) for the numeric axis. However, the other axis looks better with some sort of factor variable - so here we supply the year for the display, where the quantitative variable `year` has temporarily been used as a category/factor variable by being processed by the `factor` function before being used in the graphic:

```
ds <- gapminder %>% filter(year==1997)
#
ggplot(data=ds, mapping=aes(x=factor(year), y=lifeExp)) +
  geom_boxplot() +
  labs(x="Year",
       y="Life Expectancy (years)")
```




```
# Change orientation  
ggplot(data=ds, mapping=aes(x=factor(year),y=lifeExp)) +  
  geom_boxplot() +  
  coord_flip() +  
  labs(x="Year",y="Life Expectancy (years)")
```


Now we overlay points on top of the boxplot display. Note the `position=position_jitter` option to the `geom_point` puts some random horizontal jitter so that the points don't overlay each other. Note that the points has an argument `alpha=0.5` signifying a slightly transparent plot symbol. An `alpha` value of 1 means the plot symbol is opaque, and a value of 0 is completely transparent. Careful use of `alpha` in large datasets will enable the analyst to correctly perceive point density. Without using a smaller value of `alpha` the plot may be one large blob of ink - making it difficult to judge the density of points in the display.

```
ds <- gapminder %>% filter(year==1997)
#
ggplot(data=ds, mapping=aes(x=factor(year),y=lifeExp)) +
  geom_boxplot(outlier.shape = NA) +
  geom_point(alpha=0.5, position=position_jitter(width=0.25)) +
  labs(x="Year",y="Life Expectancy (years)")
```



If the dataframe has only one quantitative variable, we can make a character variable called “sample”, then this code will produce an acceptable display.

```
ds <- gapminder %>%
  filter(year==1997) %>%
  mutate(sample="Sample")
#
ggplot(data=ds, mapping=aes(x=sample,y=lifeExp)) +
  geom_boxplot(outlier.shape = NA) +
  geom_jitter(alpha=0.5, width=0.35) +
  labs(x="",y="Life Expectancy (years)")
```

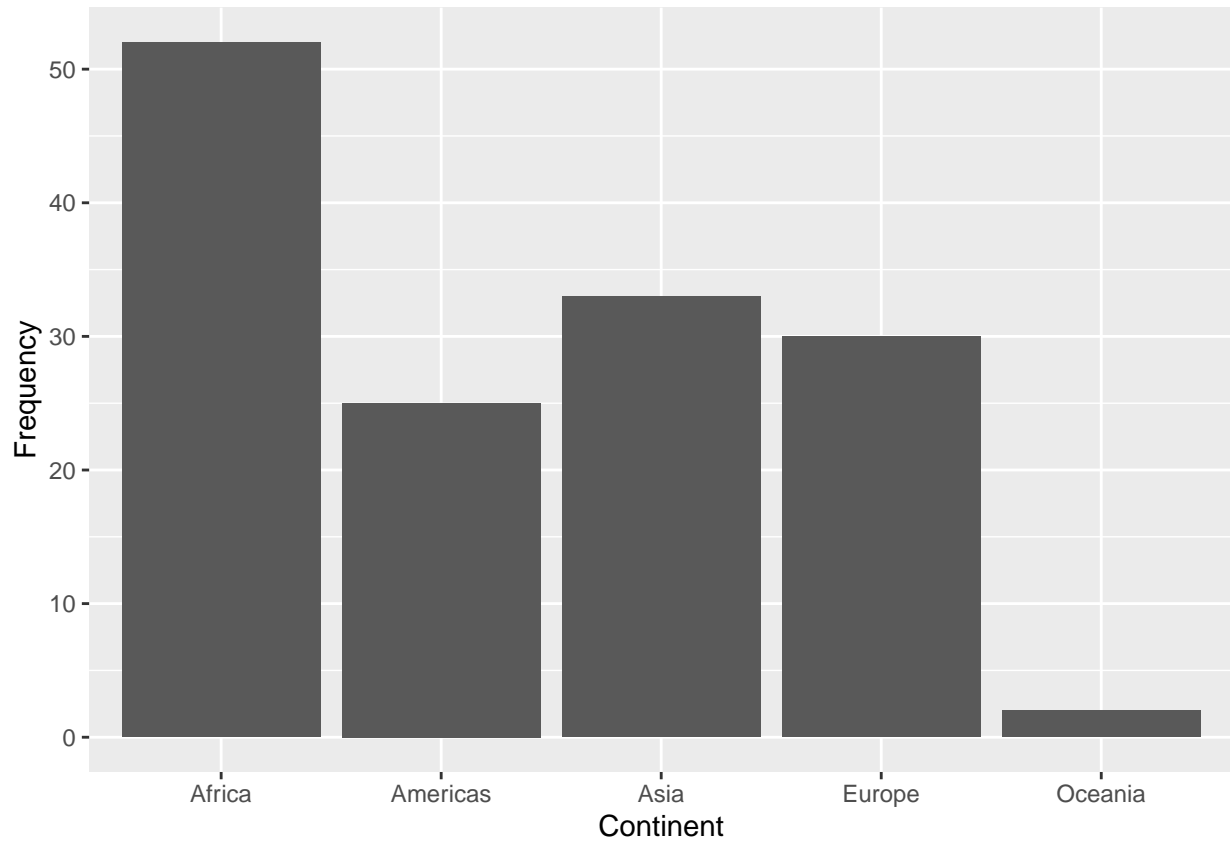


5.3 Displays of a Categorical Variable

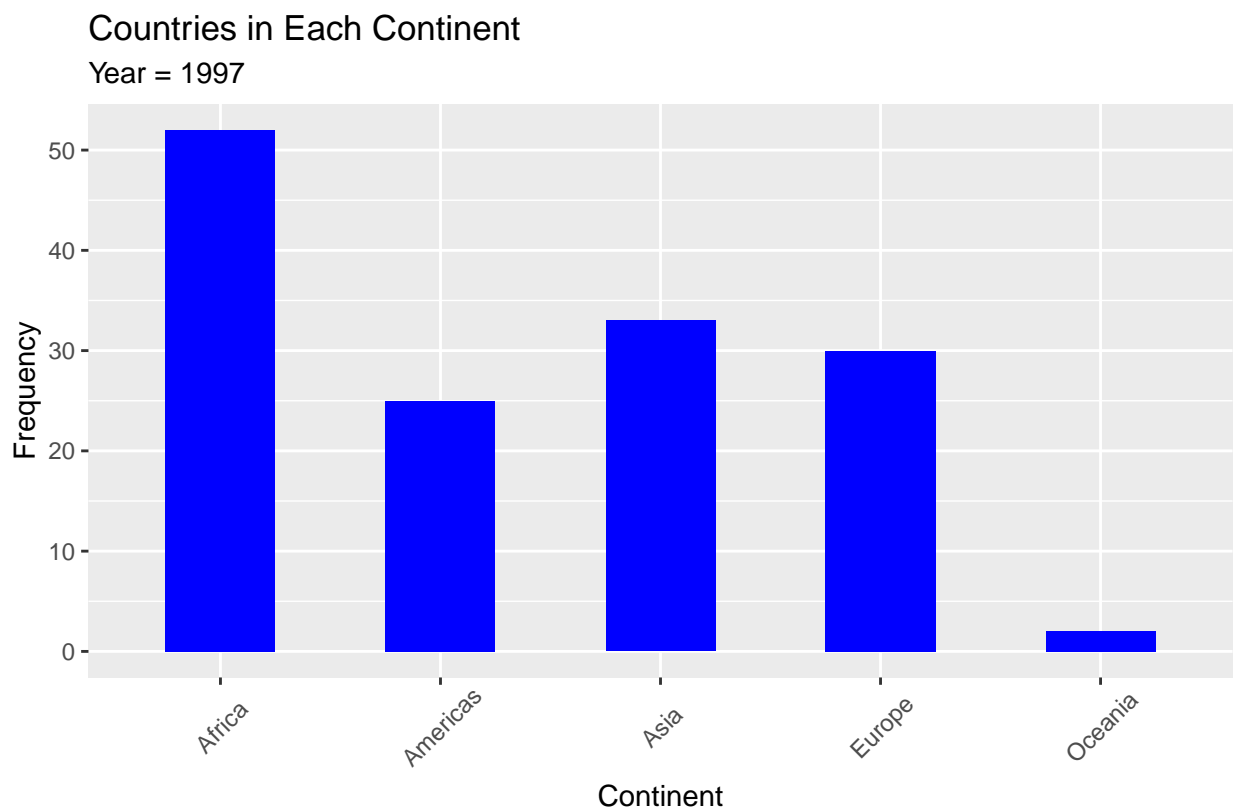
A categorical variable is summarized by using the counts for each category or the relative frequency (percentage) of each category.

5.3.1 Bar Graph

```
ds <- gapminder %>%
  filter(year==1997) %>%
  group_by(continent)
# Frequency of countries in each continent in 1997.
ggplot(data=ds, mapping=aes(x=continent)) +
  geom_bar() +
  labs(x="Continent", y="Frequency")
```




```
#  
ggplot(data=ds, mapping=aes(x=continent)) +  
  geom_bar(width=0.5, fill="blue") +  
  labs(title="Countries in Each Continent",  
        subtitle = "Year = 1997",  
        caption = "Gapminder data",  
        x="Continent",  
        y="Frequency") +  
  theme(axis.text.x = element_text(angle=45, vjust = 0.6))
```




Gapminder data

Bar graphs with percentages on vertical axis.

```
ds <- gapminder %>%  
  filter(year==1997) %>%  
  group_by(continent) %>%  
  summarise (n = n()) %>%  
  mutate(pct = 100*n / sum(n))  
#  
ggplot(data=ds, mapping=aes(x = continent, y = pct)) +  
  geom_bar(stat = "identity") +  
  xlab("Continent") + ylab("Percentage")
```




```
# change order of continents in decreasing frequency order  
ggplot(data=ds, mapping=aes(x = reorder(continent, -pct), y = pct)) +  
  geom_bar(stat = "identity") +  
  xlab("Continent") + ylab("Percentage")
```



Sometimes it is more convenient to have the bars oriented horizontally. Notice we set up the aesthetic mappings as usual and then flip the axes with the `coord_flip` command.


```
ds <- gapminder %>%
  filter(year==1997) %>%
  group_by(continent) %>%
  summarise (n = n()) %>%
  mutate(pct = 100*n / sum(n))
#
ggplot(data=ds, mapping=aes(x = reorder(continent, pct), y = pct)) +
  geom_bar(stat = "identity") +
  coord_flip() +
  xlab("Continent") + ylab("Percentage")
```



5.3.2 Pie Graph

Pie graphs are not recommended, but the code needed to make one is given here.

```
contin.prop<- gapminder %>%
  group_by(continent) %>%
  summarise (n = n()) %>%
  mutate(freq = n / sum(n))
#
ggplot(data=contin.prop, mapping=aes(x="",y=freq,fill=continent)) +
  geom_bar(width=1,stat="identity") +
  coord_polar("y",start=0) +
  xlab("") + ylab("Country Frequency by Continent")
```

Country Frequency by Continent

Chapter 6

Summary Statistics For One Variable

6.1 One Quantitative Variable

6.1.1 Using base R summary function

```
gapminder %>%  
  filter(year==1997) %>%  
  select(lifeExp) %>%  
  summary()
```

```
##      lifeExp  
##  Min.   :36.09  
## 1st Qu.:55.63  
##  Median :69.39  
##   Mean  :65.01  
## 3rd Qu.:74.17  
##   Max.  :80.69
```

6.1.2 Using dplyr summarise function

It is often helpful to create data summaries during preliminary phases of examination. Here is how to use the dplyr package summarise command in the analysis pipeline system.

```
gapminder %>%  
  dplyr::filter(year==1997) %>%  
  dplyr::summarise(meanLE=mean(lifeExp,na.rm=TRUE),  
                   medLE=median(lifeExp,na.rm=TRUE),  
                   sd=sd(lifeExp,na.rm=TRUE),  
                   iqr=IQR(lifeExp,na.rm=TRUE),  
                   Q1=quantile(lifeExp,probs=0.25,na.rm=TRUE),  
                   Q3=quantile(lifeExp,probs=0.75),  
                   n=n())
```

```
## # A tibble: 1 x 7  
##   meanLE medLE   sd   iqr   Q1   Q3     n  
##   <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <int>  
## 1   65.0  69.4  11.6  18.5  55.6  74.2   142
```

6.1.3 Summary Statistics Using gtsummary package

```
gapminder %>%
  dplyr::filter(year==1997) %>%
  dplyr::select(lifeExp,gdpPercap) %>%
  gtsummary::tbl_summary(Statistic = list(
    gtsummary::all_continuous() ~ "{mean} ({sd})",
    gtsummary::all_categorical() ~ "{n} / {N} ({p}%)",
    digits = gtsummary::all_continuous() ~ 2,
    label = c(gdpPercap ~ "Per Person GDP",
              lifeExp ~ "Life Expectancy (years)")
  )
```

```
## Table printed with `knitr::kable()`, not {gt}. Learn why at
## https://www.danielsjoberg.com/gtsummary/articles/rmarkdown.html
## To suppress this message, include `message = FALSE` in code chunk header.
```

Characteristic	**N = 142**
Life Expectancy (years)	65.01 (11.56)
Per Person GDP	9,090.18 (10,171.49)


6.1.4 Summary Statistics Using funModeling package

The `profiling_num` and `plot_num` functions from the *funModeling* package help give a concise numeric and visual overview of the numeric variables in the dataframe.

```
funModeling::profiling_num(gapminder)
```

```
##      variable      mean      std_dev variation_coef      p_01      p_05      p_25      p_
## 1      year 1.979500e+03 1.726533e+01    0.008722066    1952.0000    1952.0000    1965.750    1979.50
## 2    lifeExp 5.947444e+01 1.291711e+01    0.217187544     33.4926     38.4924     48.198     60.71
## 3      pop 2.960121e+07 1.061579e+08    3.586268548 154117.9200 475458.9000 2793664.000 7023595.50
## 4 gdpPercap 7.215327e+03 9.857455e+03    1.366182632    369.2201    547.9964    1202.060    3531.84
##      p_99      skewness      kurtosis      iqr      range_98
## 1 2.007000e+03 0.0000000    1.783217 2.750000e+01      [1952, 2007]      [19
## 2 8.023892e+01 -0.2524798    1.873099 2.264750e+01      [33.4926, 80.23892]      [41.5108
## 3 6.319900e+08 8.3328742    80.716151 1.679156e+07 [154117.92, 631990000.000002]      [946367.1, 54
## 4 3.678357e+04 3.8468819    30.431702 8.123402e+03 [369.220127794, 36783.5723707] [687.71836128, 1944
```

```
funModeling::plot_num(gapminder)
```

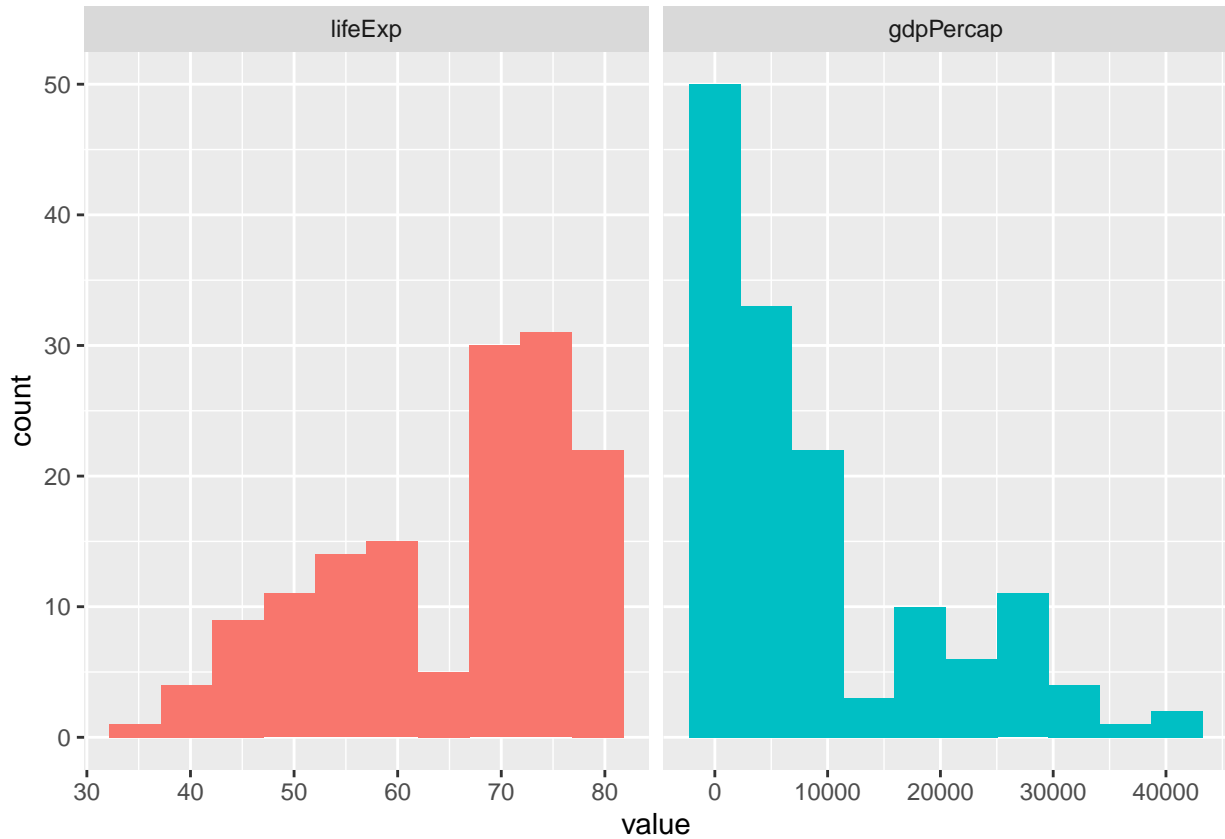


This example shows summary statistics for two quantitative variables. For only one variable, simply use `select` to obtain the desired variable.

```
gapminder %>%
  filter(year==1997) %>%
  select(lifeExp,gdpPercap) %>%
  funModeling::profiling_num()
```

```
##   variable      mean    std_dev variation_coef      p_01      p_05      p_25      p_50      p_75
## 1  lifeExp    65.01468    11.55944      0.1777974  40.03681  43.83415  55.63375  69.394   74.16975
## 2  gdpPercap 9090.17536 10171.49326      1.1189546 434.72721 590.90598 1366.83796 4781.825 12022.86719 2
##   skewness kurtosis      iqr
## 1 -0.6427906 2.218599    18.536
## 2  1.2979366 3.604446 10656.029 [40.03681, 79.7499] [47.4671, 77.548]
```

```
#
gapminder %>%
  filter(year==1997) %>%
  select(lifeExp,gdpPercap) %>%
  funModeling::plot_num()
```



6.1.5 Summary Statistics: skimr package

The *skimr* package produces summary statistics about variables and overviews for dataframes. It is easy to manipulate and use pipes, select, and filter from the *tidyverse* family of packages.

The next code supplies a dataframe that contains both categorical variables (continent), and numeric variables (lifeExp, gdpPercap). Numeric variables are chosen with the `yank` function, then some attributes are omitted from the display (`n_missing`, `complete_rate`) using the `select` function from `dplyr`.

```
varlist <- c("n_missing", "complete_rate")
gapminder %>%
  filter(year==1997) %>%
  select(-year, -country, -pop) %>%
  skimr::skim_without_charts() %>%
  skimr::yank("numeric") %>%
  dplyr::select(-one_of(varlist))
```

Variable type: numeric

skim_variable	mean	sd	p0	p25	p50	p75	p100
lifeExp	65.01	11.56	36.09	55.63	69.39	74.17	80.69
gdpPercap	9090.18	10171.49	312.19	1366.84	4781.83	12022.87	41283.16

6.2 One Categorical Variable

6.2.1 Counting Values

The next command counts the number of rows in the dataset for each continent - then we show a variant which pipes the output into the `kable` function for a more attractive table.

```
gapminder %>% count(continent)
```

```
## # A tibble: 5 x 2
##   continent     n
##   <fct>       <int>
## 1 Africa      624
## 2 Americas    300
## 3 Asia        396
## 4 Europe      360
## 5 Oceania     24
```

```
#
gapminder %>% count(continent) %>% knitr::kable()
```

continent	n
Africa	624
Americas	300
Asia	396
Europe	360
Oceania	24

```
#
gapminder %>% count(continent, sort=TRUE) %>% knitr::kable()
```

continent	n
Africa	624
Asia	396
Europe	360
Americas	300
Oceania	24

The previous code tells us how many lines (rows) for each continent, but many rows are repeated for each country - just different years. The next code counts the number of observations in variable `n` and distinct countries in variable `n_countries`.

```
gapminder %>% filter(year==1997 | year==1967) %>%
  dplyr::group_by(continent) %>%
  dplyr::summarise(n = n(), n_countries = n_distinct(country)) %>% knitr::kable()
```

continent	n	n_countries
Africa	104	52
Americas	50	25
Asia	66	33
Europe	60	30
Oceania	4	2

6.2.2 Categorical variable: gtsummary package

```
gapminder %>%
  filter(year==1997) %>%
```

```
select(continent,year) %>%
gtsummary::tbl_summary(include = c(continent),
                        statistic = list(
gtsummary::all_categorical() ~ "{n} / {N} ({p}%)")
)
```

```
## Table printed with `knitr::kable()`, not {gt}. Learn why at
## https://www.danielsjoberg.com/gtsummary/articles/rmarkdown.html
## To suppress this message, include `message = FALSE` in code chunk header.
```

Characteristic	**N = 142**
continent	
Africa	52 / 142 (37%)
Americas	25 / 142 (18%)
Asia	33 / 142 (23%)
Europe	30 / 142 (21%)
Oceania	2 / 142 (1.4%)

6.2.3 Categorical variable: skimr package

Here we summarize a categorical(factor) variable (continent), and observe it has 5 unique values (levels) and the most frequent values are displayed.

```
gapminder %>% filter(year==1997) %>%
  select(lifeExp,continent) %>%
  skimr::skim_without_charts() %>%
  skimr::yank("factor") %>%
  dplyr::select(-n_missing,-ordered,-complete_rate)
```


Variable type: factor

skim_variable	n_unique	top_counts
continent	5	Afr: 52, Asi: 33, Eur: 30, Ame: 25

6.2.4 Categorical variable: funModeling package

The *funModeling* package gives an easy way to learn about categorical variables of types: character and factor. There are two categorical variables in the gapminder dataframe: country and continent. There are a lot of countries, so we demonstrate this command for only the continent variable.

```
# Frequency distribution of entire dataframe
# will produce lots of output and warnings
funModeling::freq(gapminder)
# next command for one category variable: continent
funModeling::freq(gapminder$continent)
```

```
##      var frequency percentage cumulative_perc
## 1  Africa      624      36.62           36.62
## 2   Asia      396      23.24           59.86
## 3  Europe      360      21.13           80.99
## 4 Americas      300      17.61           98.60
## 5 Oceania       24       1.41          100.00
```

There are a lot of observations (rows) for Africa and very few for Oceania (Australia, New Zealand, etc).

6.2.5 Categorical variable: janitor package

Let's begin with the base R function `table`:

```
gapminder %>%
  filter(year==1997) %>%
  select(continent) %>%
  table()
```

```
## continent
## Africa Americas Asia Europe Oceania
##      52      25      33      30       2
```

Now contrast with the `tabyl` function from the *janitor* package:

```
gapminder %>%
  filter(year==1997) %>%
  janitor::tabyl(continent, sort=TRUE) %>%
  knitr::kable()
```

continent	n	percent
Africa	52	0.3661972
Americas	25	0.1760563
Asia	33	0.2323944
Europe	30	0.2112676
Oceania	2	0.0140845

```
#
gapminder %>%
  filter(year==1997) %>%
  janitor::tabyl(continent, sort=TRUE) %>%
  janitor::adorn_pct_formatting(digits=2, affix_sign = TRUE) %>%
  knitr::kable()
```

continent	n	percent
Africa	52	36.62%
Americas	25	17.61%
Asia	33	23.24%
Europe	30	21.13%
Oceania	2	1.41%

Chapter 7

Exploratory Data Analysis For One Quantitative Variable: by Groups

It is often helpful to create data summaries of a quantitative variable for each level of a grouping variable.

7.1 Summary Statistics: dplyr

Using *dplyr* and *tidyverse* for summary statistics across the levels of a group variable (of type factor/categorical) requires the use of the verb `group_by`. Here we produce summary statistics of life expectancy across the levels of continent.

```
# Output presented in initial continent order (alphabetic)
gapminder %>% filter(year==1997) %>%
  filter(continent != "Oceania") %>%
  group_by(continent) %>%
  summarise(meanLE=mean(lifeExp,na.rm=TRUE),
            medLE=median(lifeExp,na.rm=TRUE),
            sd=sd(lifeExp,na.rm=TRUE),
            iqr=IQR(lifeExp,na.rm=TRUE),
            Q1=quantile(lifeExp, probs=0.25,na.rm=TRUE),
            Q3=quantile(lifeExp,probs=0.75),
            n=n())
```

```
## # A tibble: 4 x 8
##   continent meanLE medLE    sd  iqr    Q1    Q3    n
##   <fct>      <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <int>
## 1 Africa      53.6  52.8  9.10 11.9  47.3  59.2   52
## 2 Americas    71.2  72.1  4.89  4.83  69.4  74.2   25
## 3 Asia        68.0  70.3  8.09 10.7  61.8  72.5   33
## 4 Europe      75.5  76.1  3.10  4.97  73.0  78.0   30
```

```
#
# Output rows ordered by decreasing values of a statistic (mean Life Expectancy):
gapminder %>% filter(year==1997) %>%
  filter(continent != "Oceania") %>%
  group_by(continent) %>%
  summarise(meanLE=mean(lifeExp,na.rm=TRUE),
            medLE=median(lifeExp,na.rm=TRUE),
            sd=sd(lifeExp,na.rm=TRUE),
```

```

    iqr=IQR(lifeExp,na.rm=TRUE),
    min=min(lifeExp),
    max=max(lifeExp),
    n=n()) %>%
  arrange(desc(meanLE))

## # A tibble: 4 x 8
##   continent meanLE medLE   sd   iqr   min   max     n
##   <fct>      <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <int>
## 1 Europe      75.5  76.1  3.10  4.97  68.8  79.4    30
## 2 Americas    71.2  72.1  4.89  4.83  56.7  78.6    25
## 3 Asia        68.0  70.3  8.09 10.7   41.8  80.7    33
## 4 Africa      53.6  52.8  9.10 11.9   36.1  74.8    52

```

Next, we save the statistics table to an object called `statstable`, then we use the `kable` function for display.

```

statstable <- gapminder %>% filter(year==1997) %>%
  filter(continent != "Oceania") %>%
  group_by(continent) %>%
  summarise(meanLE=mean(lifeExp,na.rm=TRUE),
            medLE=median(lifeExp,na.rm=TRUE),
            sd=sd(lifeExp,na.rm=TRUE),
            iqr=IQR(lifeExp,na.rm=TRUE),
            min=min(lifeExp),
            max=max(lifeExp),
            n=n()) %>%
  arrange(desc(meanLE))
#
knitr::kable(statstable)

```

continent	meanLE	medLE	sd	iqr	min	max	n
Europe	75.50517	76.116	3.104677	4.96625	68.835	79.390	30
Americas	71.15048	72.146	4.887584	4.83500	56.671	78.610	25
Asia	68.02052	70.265	8.091171	10.68100	41.763	80.690	33
Africa	53.59827	52.759	9.103387	11.92825	36.087	74.772	52

We could also “pipe” the data object into the `kable` function:

```

gapminder %>% filter(year==1997) %>%
  filter(continent != "Oceania") %>%
  group_by(continent) %>%
  summarise(meanLE=mean(lifeExp,na.rm=TRUE),
            medLE=median(lifeExp,na.rm=TRUE),
            sd=sd(lifeExp,na.rm=TRUE),
            iqr=IQR(lifeExp,na.rm=TRUE),
            min=min(lifeExp),
            max=max(lifeExp),
            n=n(),.groups = 'drop') %>%
  arrange(desc(meanLE)) %>%
  knitr::kable()

```

continent	meanLE	medLE	sd	iqr	min	max	n
Europe	75.50517	76.116	3.104677	4.96625	68.835	79.390	30
Americas	71.15048	72.146	4.887584	4.83500	56.671	78.610	25
Asia	68.02052	70.265	8.091171	10.68100	41.763	80.690	33
Africa	53.59827	52.759	9.103387	11.92825	36.087	74.772	52

7.2 Summary Statistics: gtsummary

```
#
gapminder %>%
  dplyr::filter(year==1997) %>%
  gtsummary::tbl_summary(by = continent, include = c(lifeExp, gdpPercap),
    statistic = list(
      gtsummary::all_continuous() ~ "{mean} ({sd})",
      gtsummary::all_categorical() ~ "{n} / {N} ({p}%)" ),
  digits = gtsummary::all_continuous() ~ 2,
  label = c(gdpPercap ~ "Per Person GDP",
    continent ~ "Continent",
    lifeExp ~ "Life Expectancy (years)")
  )
```

Table printed with `knitr::kable()`, not {gt}. Learn why at
 ## <https://www.danielsjoberg.com/gtsummary/articles/rmarkdown.html>
 ## To suppress this message, include `message = FALSE` in code chunk header.

Characteristic	**Africa**, N = 52	**Americas**, N = 25	**Asia**, N = 33	**Europe**, N = 30	**Oc
Life Expectancy (years)	53.60 (9.10)	71.15 (4.89)	68.02 (8.09)	75.51 (3.10)	7
Per Person GDP	2,378.76 (2,820.73)	8,889.30 (7,874.23)	9,834.09 (11,094.18)	19,076.78 (10,065.46)	24,02

7.3 Summary Statistics: skimr

Here we implement the `group_by` function to display descriptive statistics for numeric variables by continent, for two quantitative variables using functions from the *skimr* package.

```
gapminder %>% filter(year==1997) %>%
  filter(continent != "Oceania") %>%
  group_by(continent) %>%
  skimr::skim_without_charts() %>%
  skimr::yank("numeric") %>%
  dplyr::filter(skim_variable %in% c("lifeExp", "gdpPercap")) %>%
  knitr::kable()
```


skim_variable	continent	n_missing	complete_rate	mean	sd	p0	p25	p50
lifeExp	Africa	0	1	53.59827	9.103387	36.0870	47.30025	52.75
lifeExp	Americas	0	1	71.15048	4.887584	56.6710	69.38800	72.14
lifeExp	Asia	0	1	68.02052	8.091171	41.7630	61.81800	70.26
lifeExp	Europe	0	1	75.50517	3.104677	68.8350	73.02350	76.11
gdpPercap	Africa	0	1	2378.75956	2820.728117	312.1884	791.90197	1179.88
gdpPercap	Americas	0	1	8889.30086	7874.225145	1341.7269	4684.31381	7113.69
gdpPercap	Asia	0	1	9834.09330	11094.180481	415.0000	1902.25210	3645.38
gdpPercap	Europe	0	1	19076.78180	10065.457716	3193.0546	9946.59931	19596.49

7.4 Graphical Displays of a quantitative variable, separated by groups

In each example, the first lines create the dataset to be graphed - followed by a `ggplot` command making the display. Several of the examples make use of the principle of “small-multiples” so that each level of the factor variable has a separate panel for the quantitative variable display.

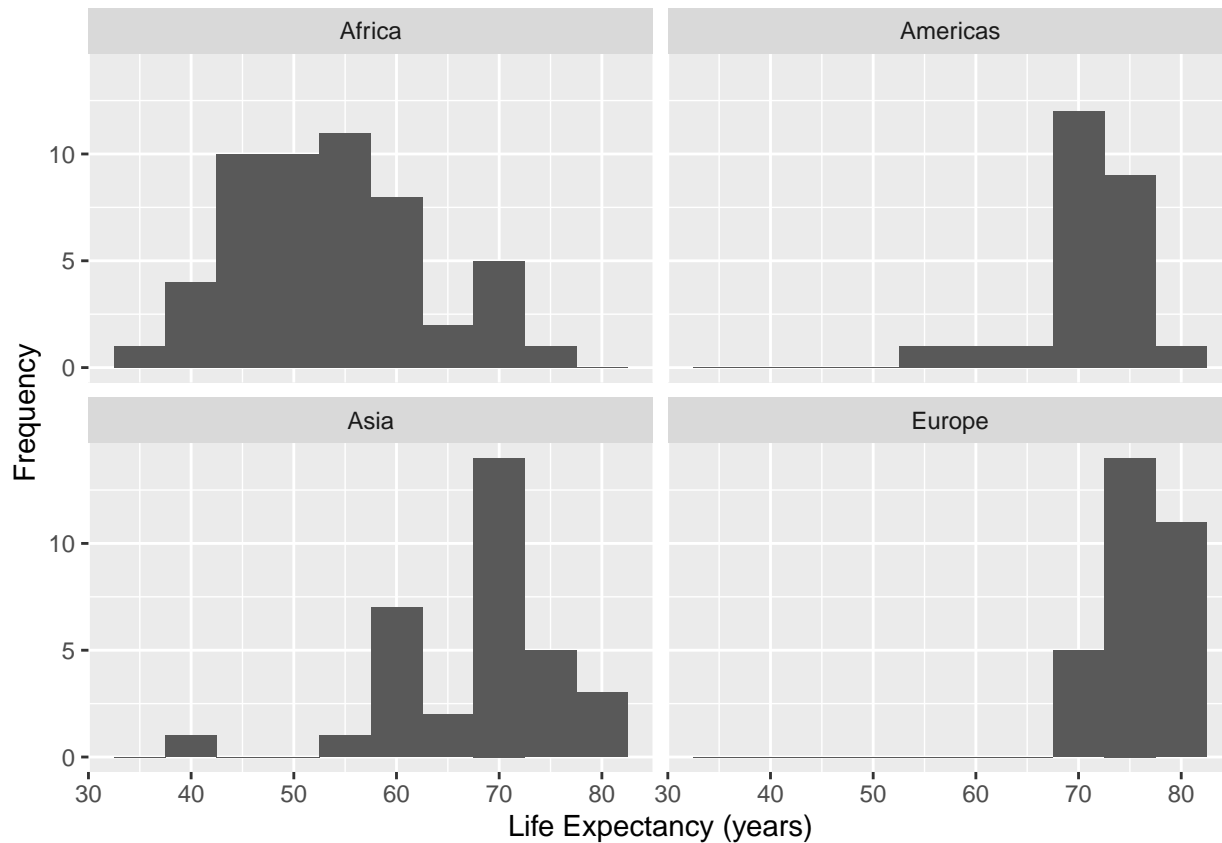
7.4.1 Dotplots

```
ds <- gapminder %>% filter(year==1997)
#
ggplot(data=ds, mapping=aes(x=lifeExp)) +
  geom_dotplot() +
  facet_wrap( ~ continent, ncol=2) +
  xlab("Life Expectancy (years)") +
  ylab("Frequency")
```



7.4.2 Histograms

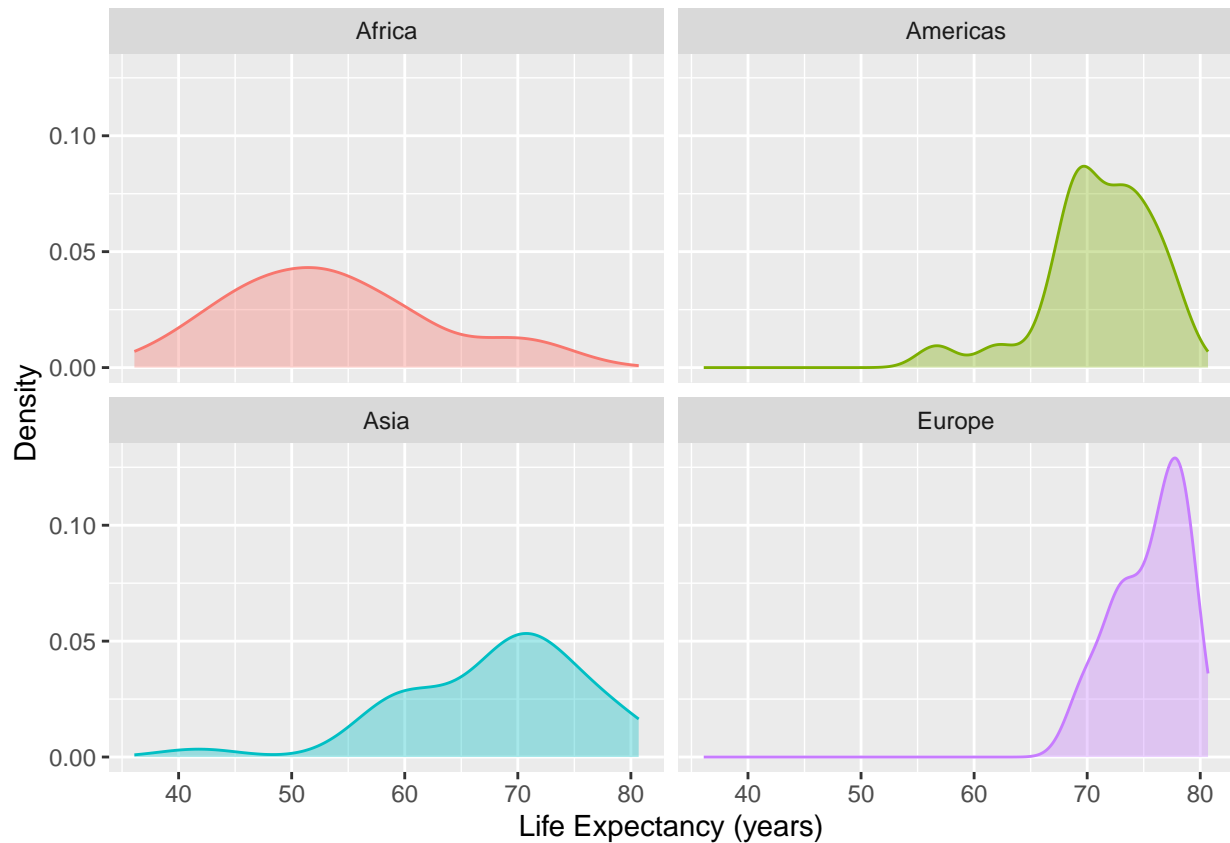
```
gapminder %>%
  filter(year==1997) %>%
  filter(continent != "Oceania") %>%
  group_by(continent) %>%
  ggplot(data=., mapping=aes(x=lifeExp)) +
  geom_histogram(binwidth=5) +
  facet_wrap( ~ continent, ncol=2) +
  xlab("Life Expectancy (years)") +
  ylab("Frequency")
```



7.4.3 Density Plots in Facets

The code given here shows how to produce a density plot in separate panels for each continent.


```
gapminder %>%
  filter(year==1997) %>%
  filter(continent != "Oceania") %>%
  group_by(continent) %>%
  ggplot(data=., mapping=aes(x=lifeExp, colour=continent, fill=continent)) +
  geom_density(alpha = 0.35) +
  xlab("Life Expectancy (years)") +
  ylab("Density") +
  facet_wrap( ~ continent, ncol = 2) +
  theme(legend.position = "none")
```



7.4.4 Overlaid Density Plots

The initial command below takes the `gapminder` data and consider only observations (rows) from 1997, but exclude all observations from Oceania. The `alpha` setting controls the amount of transparency in the densities for each continent - smaller values of `alpha` (between 0 and 1) are more transparent.

```
gapminder %>%
  filter(year==1997) %>%
  filter(continent != "Oceania") %>%
  group_by(continent) %>%
  ggplot(data=., mapping=aes(x=lifeExp, colour=continent, fill=continent)) +
  geom_density(alpha = 0.35) +
  xlab("Life Expectancy (years)") +
  ylab("Density")
```





7.4.5 Boxplots, Grouped Data


In the code below, the `alpha` value again controls the transparency of the points `alpha=1` means opaque, `alpha=0` means completely see-through. When there is a lot of data, use a smaller value of `alpha`.

The last set of boxplots uses both vertical(height) and horizontal(width) jitter to prevent points being overlaid, and be more visible.

```
ds <- gapminder %>%
  filter(year==1997) %>%
  filter(continent != "Oceania") %>%
  group_by(continent)
#
ggplot(data=ds, mapping=aes(x=continent,y=lifeExp)) +
  geom_boxplot() +
  labs(x="Continent",y="Life Expectancy (years)")
```




```
#
ggplot(data=ds, mapping=aes(x=continent,y=lifeExp)) +
  geom_boxplot(outlier.colour = NA) +
  geom_point(position = position_jitter(width = 0.15, height = 0.15),alpha=.50) +
  labs(x="Continent",y="Life Expectancy (years)")
```




7.4.6 Boxplots, overlay points on the boxplots with color control

In the code below, the alpha value controls the transparency of the points alpha=1 means opaque, alpha=0 means completely see-through.

```
ds <- gapminder %>%
  filter(year==1997) %>%
  filter(continent != "Oceania") %>%
  group_by(continent)
#
ggplot(data=ds, mapping=aes(x=continent,y=lifeExp, colour=continent)) +
  geom_point(position = position_jitter(width = 0.2, height = 0.2),alpha=.25) +
  geom_boxplot(outlier.colour = NA, fill = NA) +
  labs(x="Continent",y="Life Expectancy (years)")
```



```
#
ggplot(data=ds, mapping=aes(x=continent,y=lifeExp, colour=continent)) +
  geom_point(position = position_jitter(width = 0.2, height = 0.2),alpha=.80) +
  geom_boxplot(outlier.colour = NA, fill = NA) +
  labs(x="Continent",y="Life Expectancy (years)")
```



Chapter 8

Analysis of One Categorical Variable by another categorical variable

To demonstrate graphical displays of two categorical variables, we need a new dataset with two categorical variables. We use the `congress_age` dataframe from the *fivethirtyeight* package. In these displays we will use categorical variables:

- party affiliation (party) with values: D, I, R.
- congressional chamber (chamber) with values: house, senate

We will restrict ourselves to the 113th congress, a meeting of the legislative branch of the United States federal government, from January 3, 2013, to January 3, 2015, during the fifth and sixth years of Barack Obama's presidency.

8.1 Tables

8.1.1 Base Functions

```
congage <- fivethirtyeight::congress_age
ds1 <- congage %>% filter(congress > 112) %>% select(congress, chamber, state, party, incumbent, age)
# We declare party and chamber as factor/categorical variables, and control their levels.
ds1 <- ds1 %>% mutate(party = factor(party, levels=c("D", "I", "R")),
                     chamber = factor(chamber))
ds1 <- ds1 %>% na.omit()
ds <- ds1
#
table(ds$chamber, ds$party)

##
##           D    I    R
## house  202    0  237
## senate   57    2   46
#
mytable <- table(ds$chamber, ds$party)
#
prop.table(mytable) # cell percentages

##
```

```
##           D           I           R
##  house 0.371323529 0.000000000 0.435661765
##  senate 0.104779412 0.003676471 0.084558824
```

```
prop.table(mytable, 1) # row percentages
```

```
##
##           D           I           R
##  house 0.46013667 0.00000000 0.53986333
##  senate 0.54285714 0.01904762 0.43809524
```

```
prop.table(mytable, 2) # column percentages
```

```
##
##           D           I           R
##  house 0.7799228 0.0000000 0.8374558
##  senate 0.2200772 1.0000000 0.1625442
```

8.1.2 janitor package tools

```
ds %>% janitor::tabyl(chamber, party)
```

```
## chamber   D I   R
##   house 202 0 237
##   senate  57 2  46
```

```
#
t2 <- ds %>% janitor::tabyl(chamber, party)
t2 %>%
  janitor::adorn_percentages("row") %>%
  janitor::adorn_pct_formatting(digits = 2) %>%
  janitor::adorn_ns()
```

```
## chamber           D           I           R
##   house 46.01% (202) 0.00% (0) 53.99% (237)
##   senate 54.29% (57) 1.90% (2) 43.81% (46)
```

```
# column percentages
```

```
t2 %>%
  janitor::adorn_percentages("col") %>%
  janitor::adorn_pct_formatting(digits = 2) %>%
  janitor::adorn_ns()
```

```
## chamber           D           I           R
##   house 77.99% (202) 0.00% (0) 83.75% (237)
##   senate 22.01% (57) 100.00% (2) 16.25% (46)
```

```
# both row and column percentages
```

```
t2 %>%
  janitor::adorn_percentages("all") %>%
  janitor::adorn_pct_formatting(digits = 2) %>%
  janitor::adorn_ns()
```

```
## chamber           D           I           R
##   house 37.13% (202) 0.00% (0) 43.57% (237)
##   senate 10.48% (57) 0.37% (2) 8.46% (46)
```


8.1.3 tidyverse tools

```

congage <- fivethirtyeight::congress_age
ds1 <- congage %>% filter(congress > 112) %>% select(congress, chamber, state, party, incumbent, age)
# We declare party and chamber as factor/categorical variables, and control their levels.
ds1 <- ds1 %>% mutate(party = factor(party, levels=c("D", "I", "R")),
                     chamber = factor(chamber))
ds1 <- ds1 %>% na.omit()
ds <- ds1
ds %>% group_by(chamber, party) %>%
  dplyr::count() %>%
  tidyr::pivot_wider(names_from = party, values_from = n)

## # A tibble: 2 x 4
## # Groups:   chamber [2]
##   chamber    D     R     I
##   <fct>   <int> <int> <int>
## 1 house    202   237    NA
## 2 senate    57    46     2

```

8.1.4 gtsummary tools

```

congage <- fivethirtyeight::congress_age
ds1 <- congage %>% dplyr::filter(congress > 112) %>% dplyr::select(congress, chamber, state, party, incumbent,
# We declare party and chamber as factor/categorical variables, and control their levels.
ds1 <- ds1 %>% mutate(party = factor(party, levels=c("D", "I", "R")),
                     chamber = factor(chamber)) %>% dplyr::select(party, chamber)
ds1 %>%
  na.omit() %>%
  gtsummary::tbl_summary(
    by = chamber,
    statistic=gtsummary::all_categorical() ~ "{n} / {N} ({p}%)"
  )

## Table printed with `knitr::kable()`, not {gt}. Learn why at
## https://www.danielsjoberg.com/gtsummary/articles/rmarkdown.html
## To suppress this message, include `message = FALSE` in code chunk header.

```


Characteristic	**house**, N = 439	**senate**, N = 105
party		
D	202 / 439 (46%)	57 / 105 (54%)
I	0 / 439 (0%)	2 / 105 (1.9%)
R	237 / 439 (54%)	46 / 105 (44%)

8.2 Graphical Displays


```

# basic bar plot of party affiliation
ggplot(data=ds, mapping=aes(x=party)) +
  geom_bar() +
  labs(x="Party", y="Count")


```



```
ds <- ds1 %>% group_by(party, chamber)
#
ggplot(data=ds, mapping=aes(x=chamber)) +
  geom_bar(aes(fill=party), position="dodge") +
  labs(x="Chamber", y="Count")
```



```
#  
ggplot(data=ds, mapping=aes(x=party)) +  
  geom_bar(aes(fill=party)) +  
  facet_wrap( ~ chamber) +  
  labs(x="Party", y="Count")
```




*# The next display attempts to use percentages on the vertical axis defined within chamber.
This means the next command must list chamber as the FIRST group_by variable.*

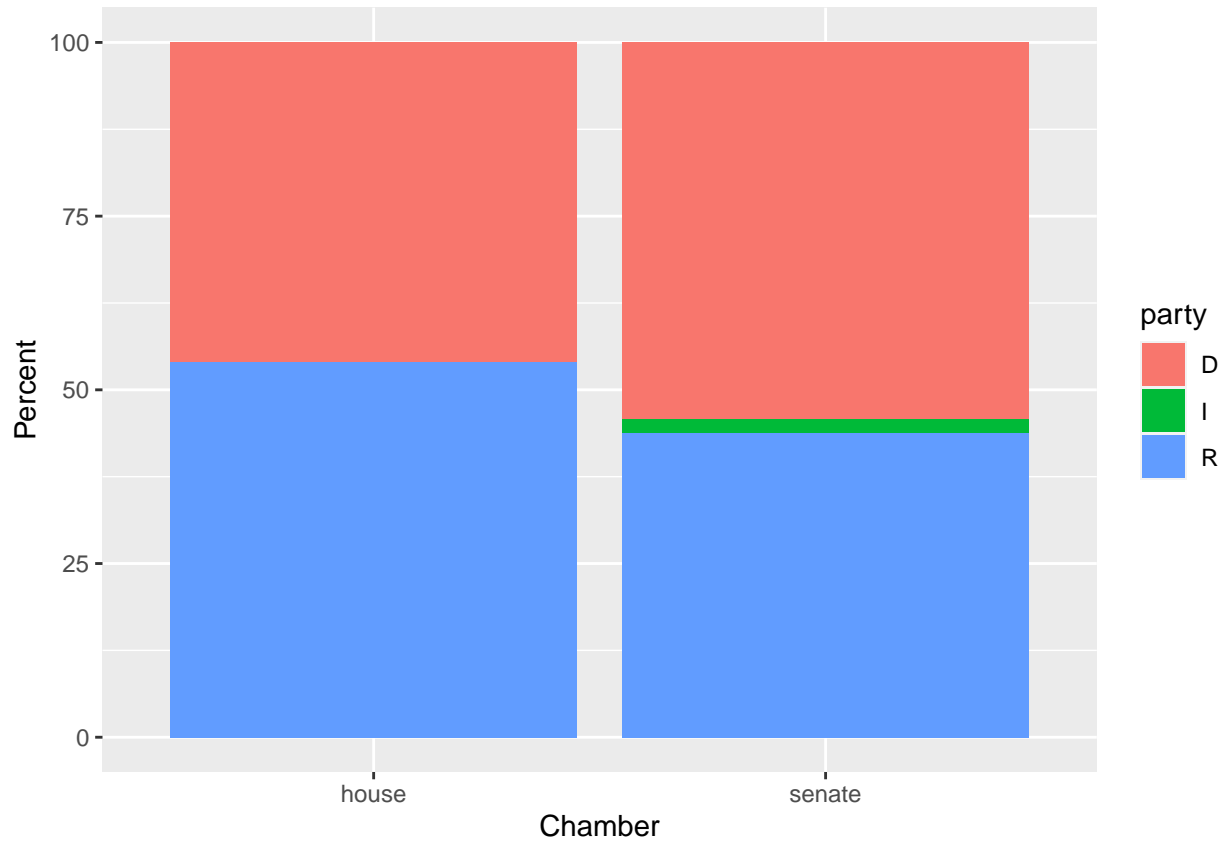
```
ds <- ds1 %>% group_by(chamber,party) %>%
  summarise (n = n()) %>%
  mutate(pct = 100*n / sum(n))
#
ds
```

```
## # A tibble: 5 x 4
## # Groups:   chamber [2]
##   chamber party    n  pct
##   <fct>   <fct> <int> <dbl>
## 1 house    D      202  46.0
## 2 house    R      237  54.0
## 3 senate   D       57  54.3
## 4 senate   I        2   1.90
## 5 senate   R       46  43.8
```

```
#
ggplot(data=ds, mapping=aes(x=party, y=pct)) +
  geom_bar(aes(fill=party),stat="identity") +
  facet_wrap( ~ chamber) +
  labs(x="Party", y="Percent")
```



```
ds1 <- congage %>% filter(congress > 112) %>% select(congress, chamber, state, party, incumbent, age) %>%
  mutate(party = factor(party, levels=c("D", "I", "R")),
         chamber = factor(chamber)) %>%
  na.omit()
#
ds <- ds1 %>% group_by(chamber, party) %>%
  summarise(n = n()) %>%
  mutate(pct = 100*n / sum(n))
#
ggplot(data=ds, mapping=aes(x=chamber, y=pct, fill=party)) +
  geom_col() +
  labs(x="Chamber", y="Percent")
```



```
ds1 <- congage %>% filter(congress > 112) %>% select(congress, chamber, state, party, incumbent, age) %>%
  mutate(party = factor(party, levels=c("D", "I", "R")),
         chamber = factor(chamber)) %>%
  na.omit()
#
ds1 %>% group_by(chamber, party) %>%
  summarise(n = n()) %>%
  mutate(pct = 100*n / sum(n)) %>%
ggplot(data=., mapping=aes(x=chamber, y=pct, fill=party)) +
  geom_col() +
  coord_flip() +
  labs(x="Chamber", y="Percent")
```

